

# A simple neural classifier

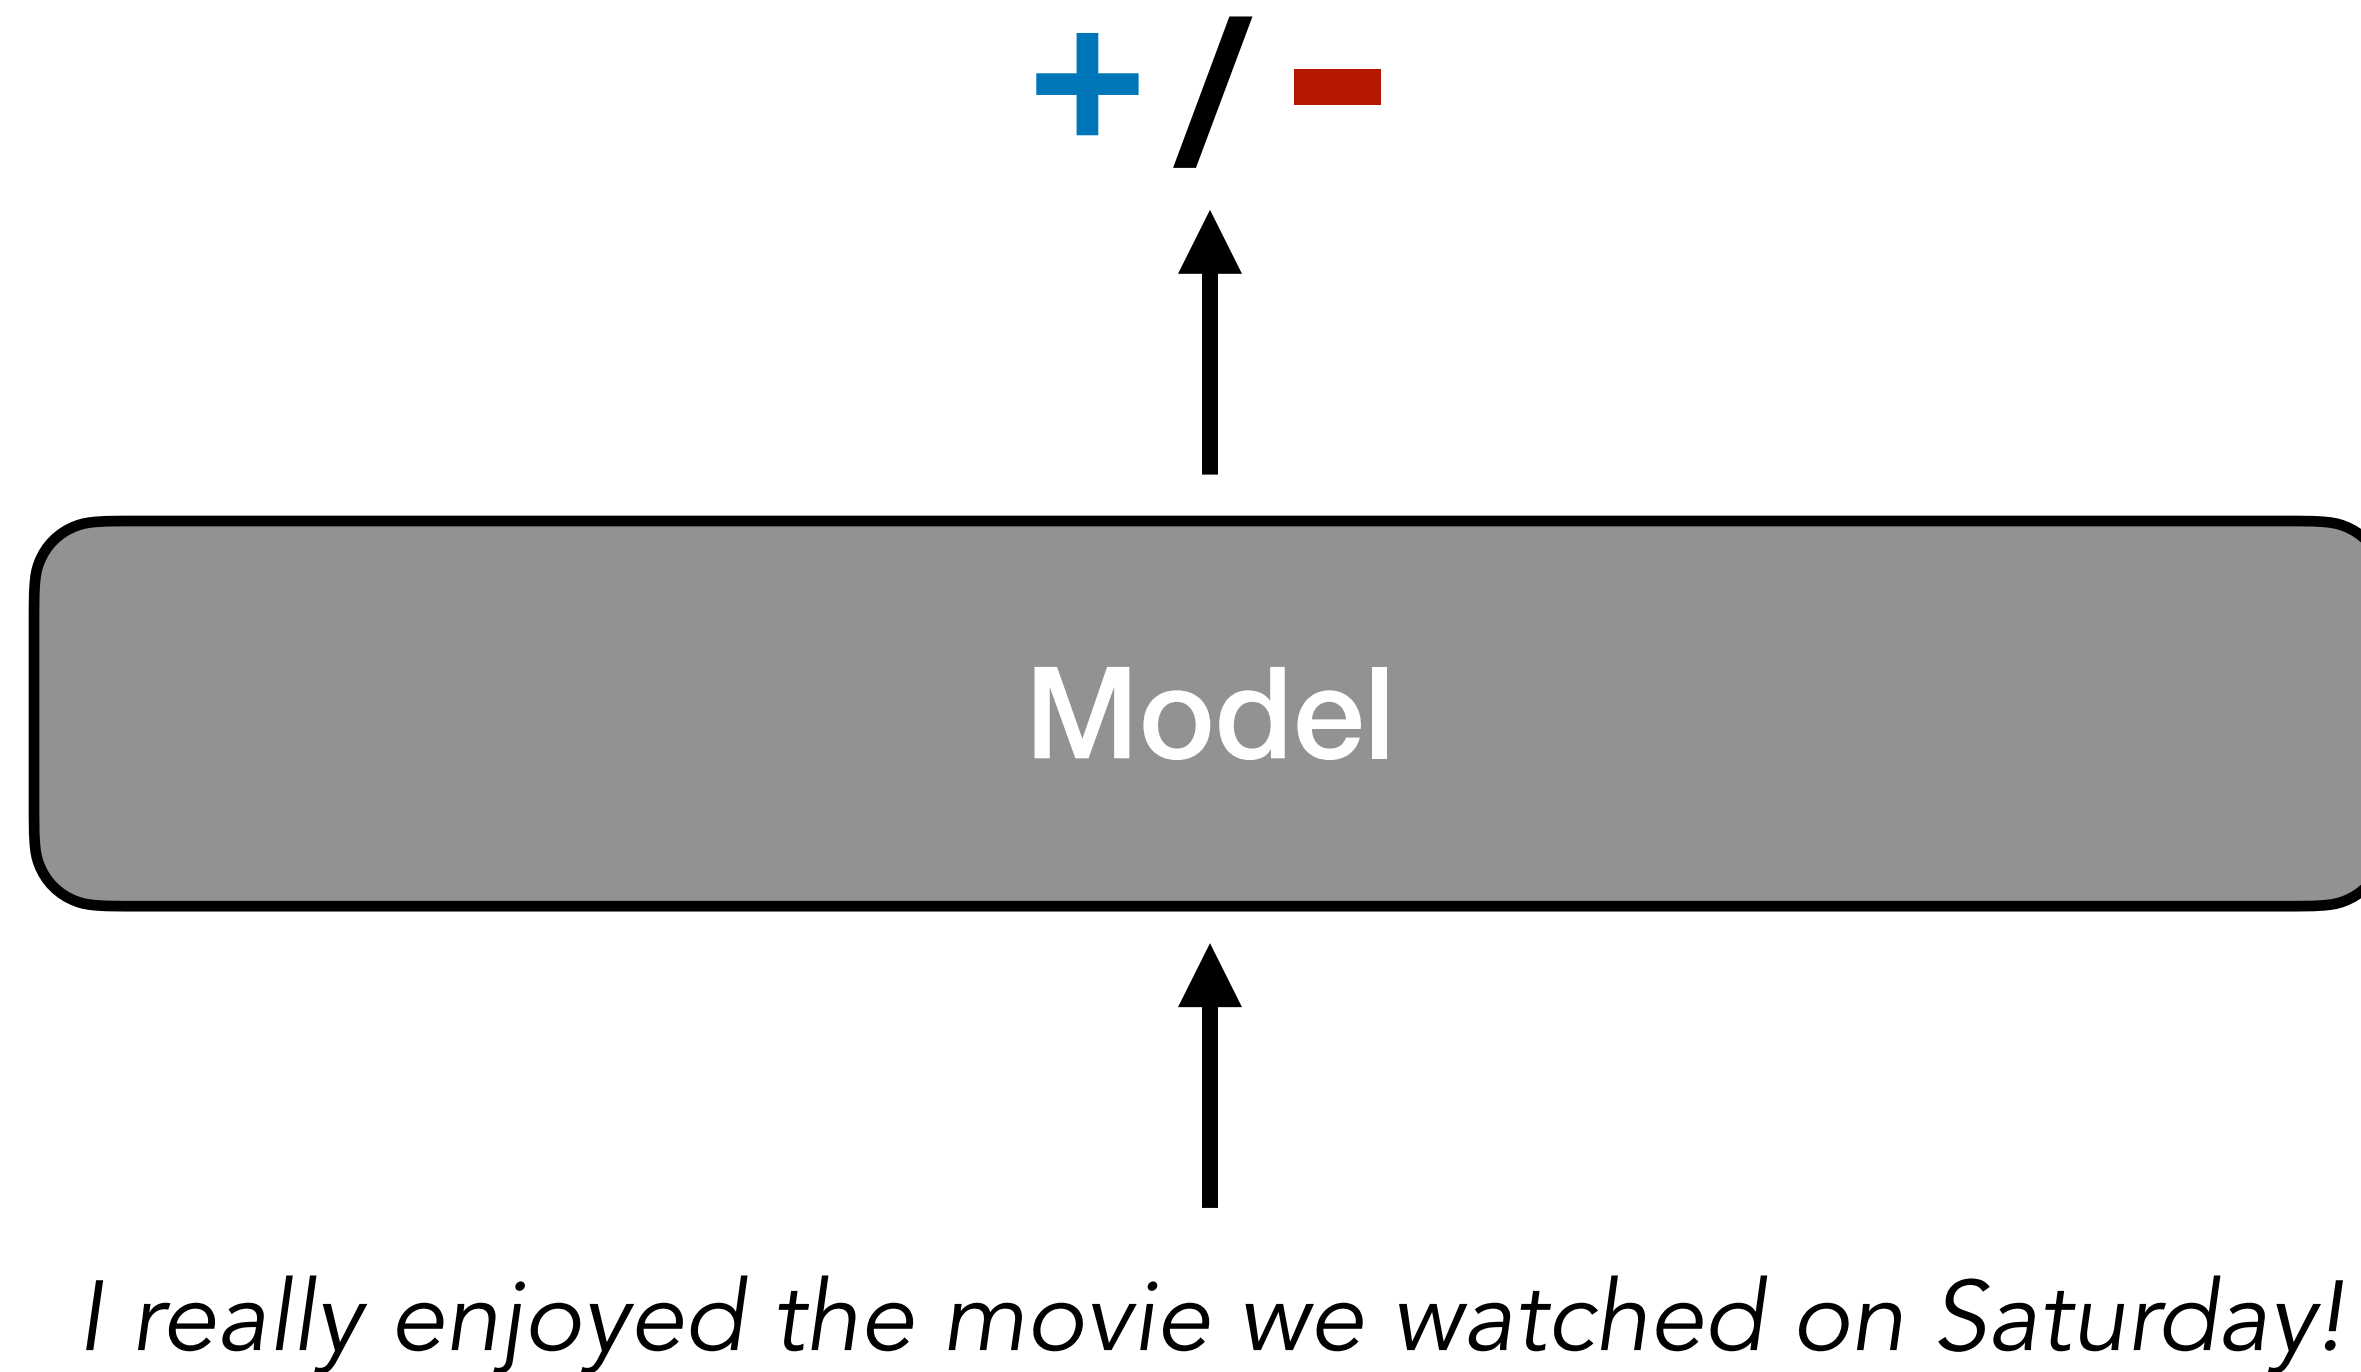
Antoine Bosselut

# Section Outline

- Setting up an NLP problem
- **Embeddings** - how do we represent sequences of discrete words ?
- **Model** - how do we compose our embeddings into higher-level representations?
- **Prediction** - how do we map our model's representation of the task to a prediction?

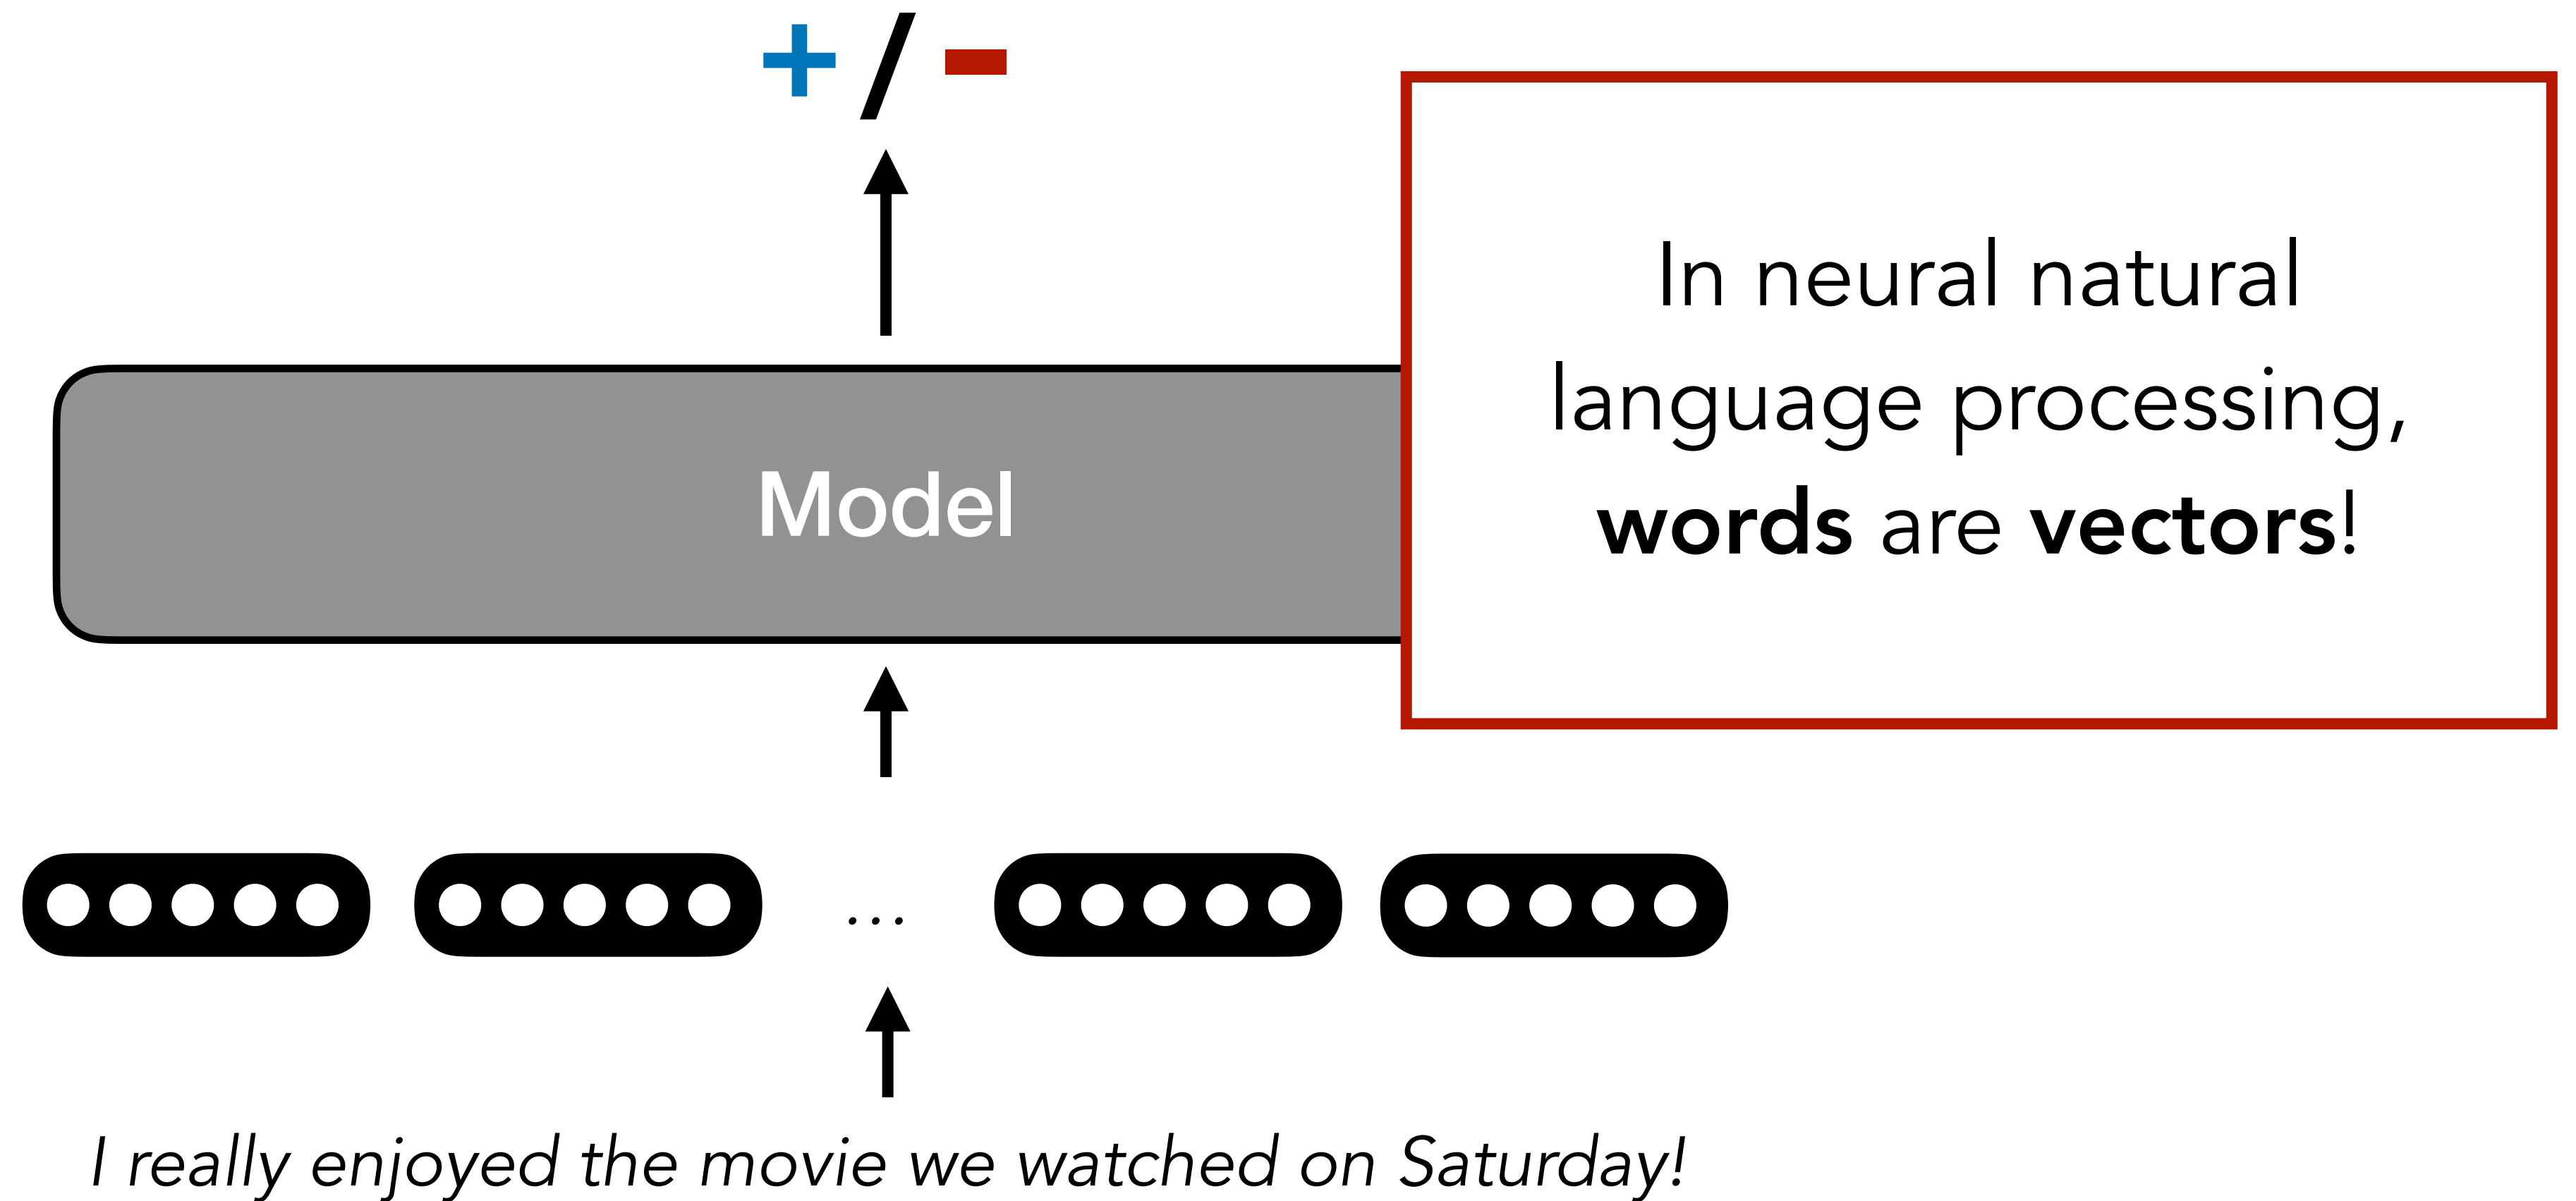
# A simple NLP model

- **Example:** Convert a sentence describing a movie review to a sentiment



# A simple NLP model

- **Example:** Convert a sentence describing a movie review to a sentiment



# Question

**What words should we model as vectors?**

# Choosing a vocabulary

- Language contains many words (e.g., ~600,000 in English)
  - **What about other tokens:** Capitalisation? Accents ? Typos!? Words in other languages!? In other scripts!? Emojis !? Unicode !?
  - **Millions of potential unique tokens!** Most rarely appear in our training data (Zipfian distribution)
  - Model has limited capacity

# Choosing a vocabulary

- Language contains many words (e.g., ~600,000 in English)
  - **What about other tokens:** Capitalisation? Accents ? Typos!? Words in other languages!? In other scripts!? Emojis !? Unicode !?
  - **Millions of potential unique tokens!** Most rarely appear in our training data (Zipfian distribution)
  - Model has limited capacity
- How should we select which tokens we want our model to process?
  - Week 11 - tokenisation!
  - For now, initialize a vocabulary  $V$  of tokens that we can represent as a vector
  - Any token not in this vocabulary  $V$  is mapped to a special <UNK> token (e.g., unknown).

# Question

**How should we model a word as a vector?**



# One upon a time: **sparse word representations**

$$x_i \in \{0,1\}^V$$

- Define a vocabulary  $V$
- Each word in the vocabulary is represented by a sparse vector
- Dimensionality of sparse vector is size of vocabulary (e.g., thousands, possibly millions)

<i>I</i>	→	<i>[ 0 ... 0 0 0 1 ... 0 0 ]</i>
<i>really</i>	→	<i>[ 0 ... 1 ... 0 0 0 0 0 ]</i>
<i>enjoyed</i>	→	<i>[ 0 ... 0 0 0 1 0 ... 0 ]</i>
<i>the</i>	→	<i>[ 0 ... 0 1 0 0 0 ... 0 ]</i>
<i>movie</i>	→	<i>[ 0 ... 0 0 0 0 0 ... 1 ]</i>
<i>!</i>	→	<i>[ 1 ... 0 0 0 0 0 0 0 0 ]</i>

# Word Vector Composition

- To represent sequences, beyond single words, define a composition function over sparse vectors

*I really enjoyed the movie !* → [ 1 ... 1 1 0 1 ... 0 1 ]      Simple Counts

*I really enjoyed the movie !* → [ 0.01 ... 0.1 0.1 0 0.001 ... 0 0.5 ]      Weighted by Corpus Statistics (e.g., TF-IDF)

Many others...

# Problem

**With sparse vectors, similarity is a function of common words!**

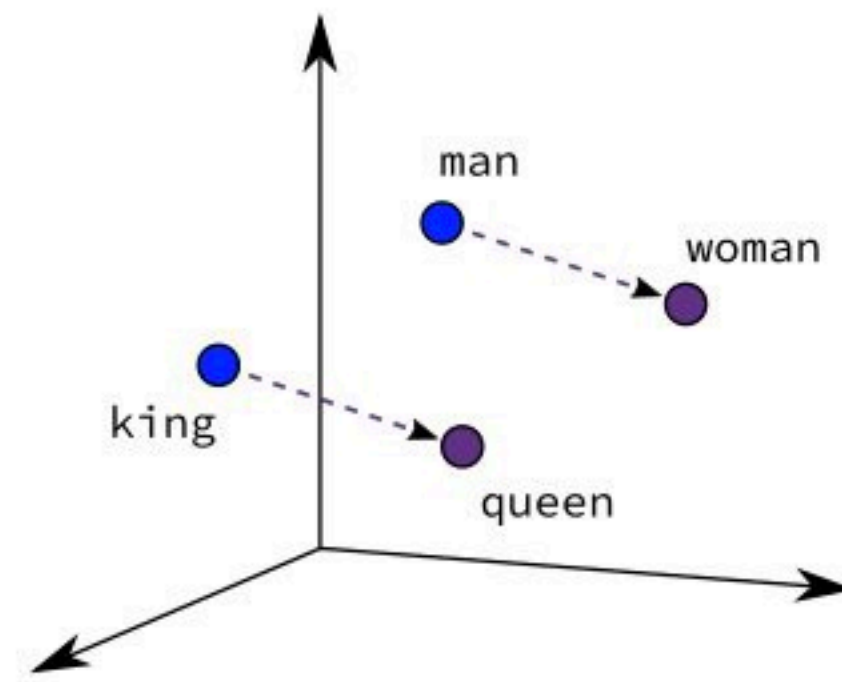
**How do you learn learn similarity between words?**

*enjoyed*  $\longrightarrow$   $[0 \dots 0 \ 0 \ 0 \ 1 \dots 0 \ 0]$

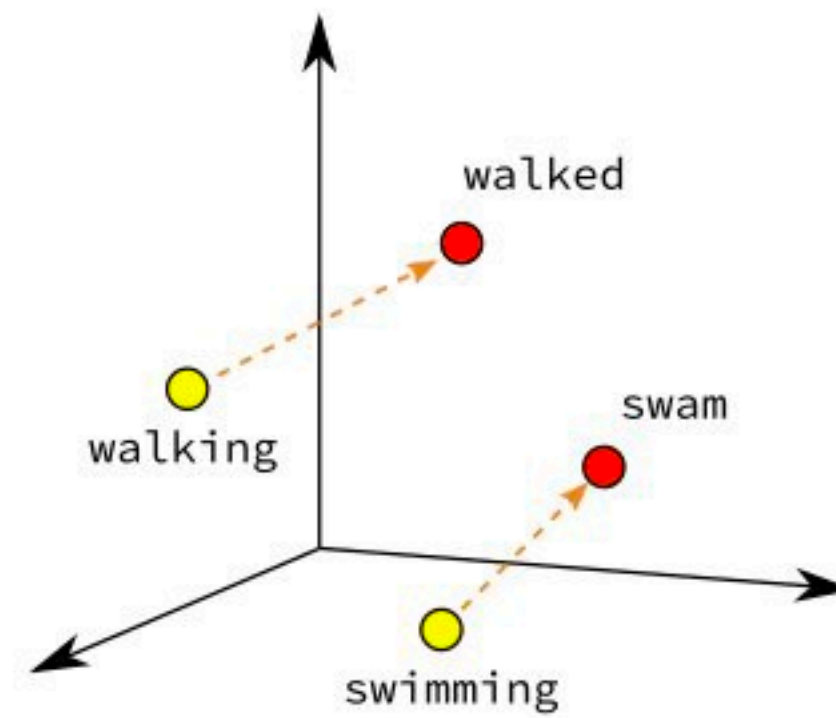
*loved*  $\longrightarrow$   $[0 \dots 1 \dots 0 \ 0 \ 0 \ 0 \ 0]$

$\text{sim}( \textit{enjoyed}, \textit{loved} ) = 0$

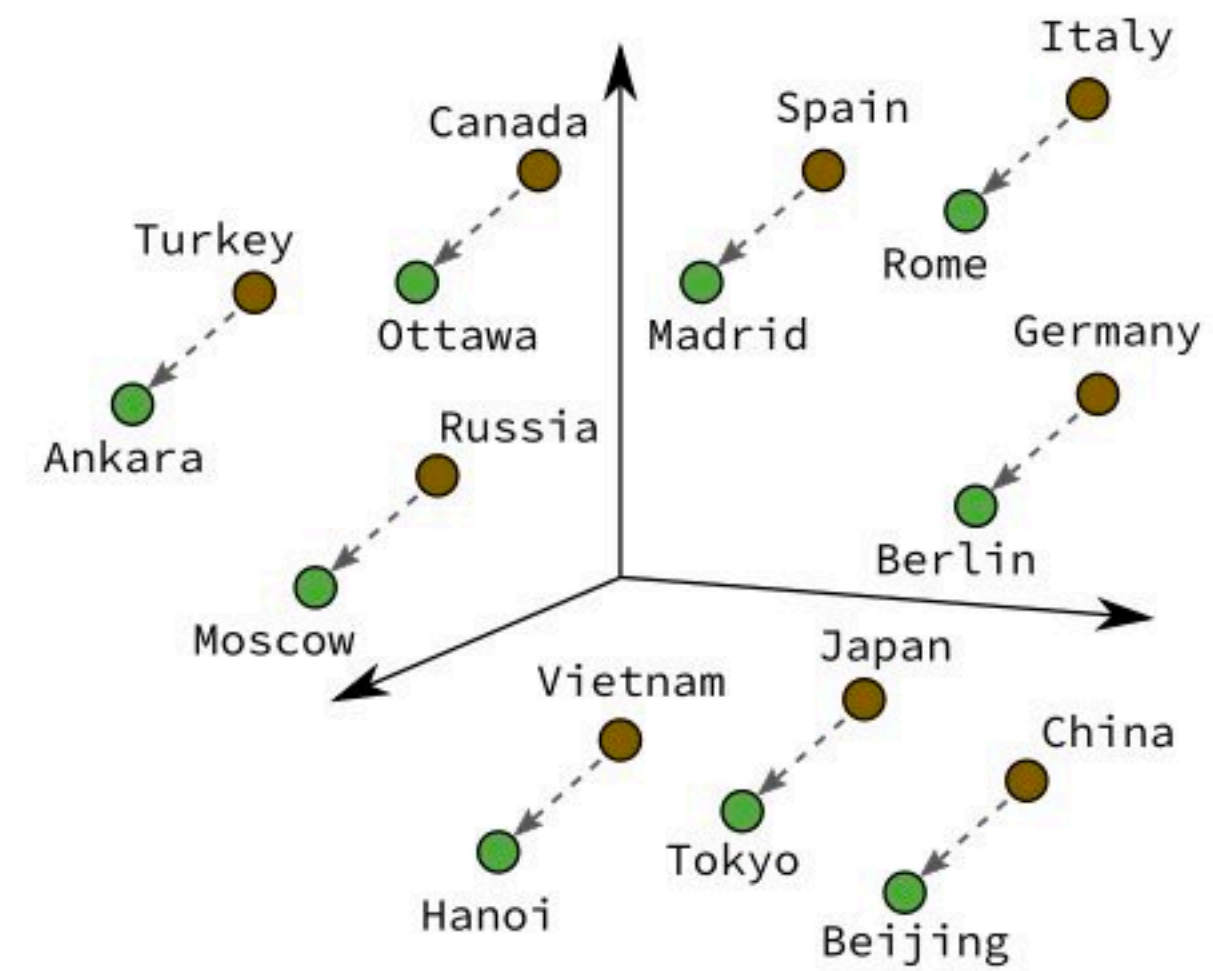
# Embeddings Goal



Male-Female



Verb Tense



Country-Capital

**How do we train semantics-encoding embeddings of words?**

# Dense Word Vectors

- Represent each word as a high-dimensional\*, **real-valued** vector
  - \*Low-dimensional compared to V-dimension sparse representations, but still usually  $O(10^2 - 10^3)$

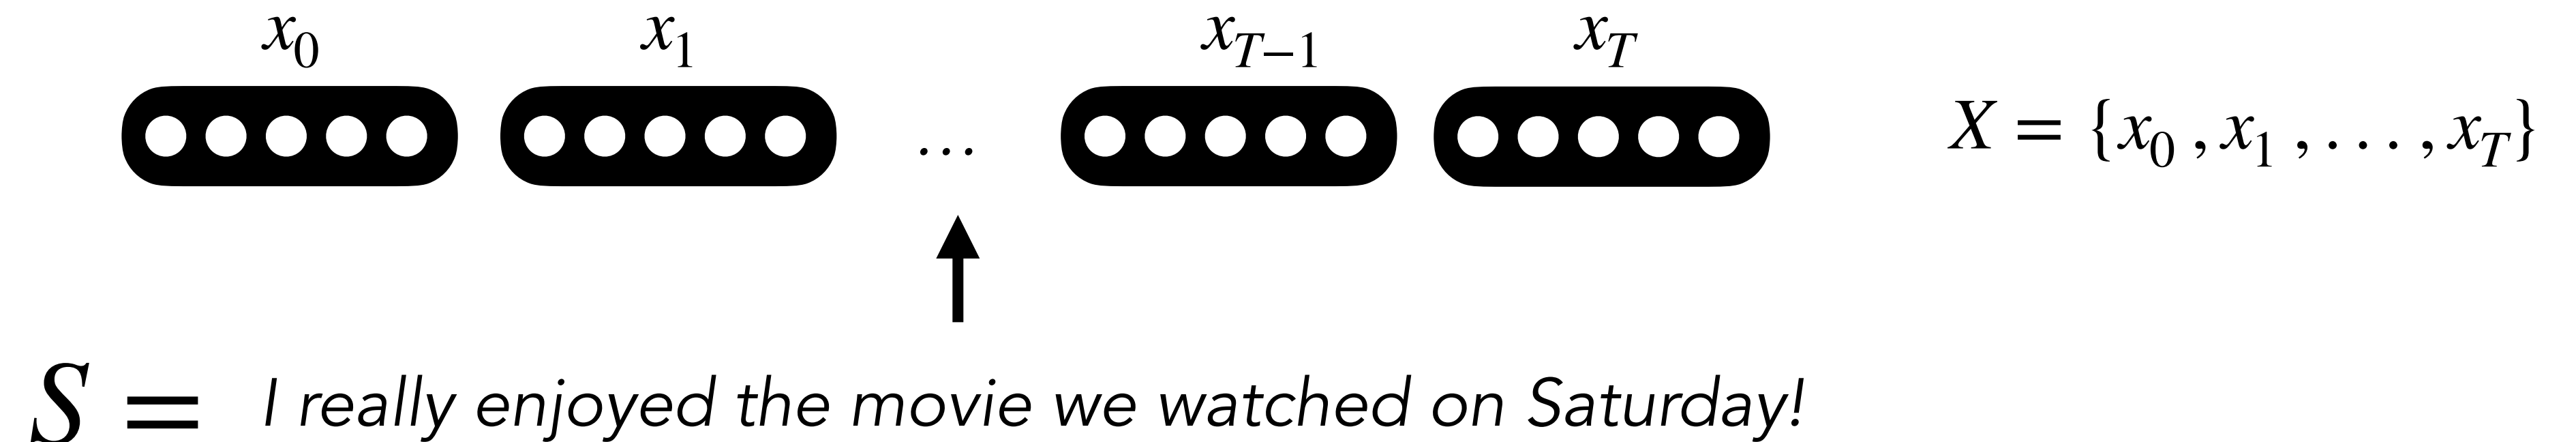
<i>I</i>	→	[ 0.113 -0.782 1.893 0.984 6.349 ... ]
<i>really</i>	→	[ 0.906 0.661 -0.214 -0.894 -0.880 ... ]
<i>enjoyed</i>	→	[ -0.842 0.647 -0.882 0.045 0.029 ... ]
<i>the</i>	→	[ 0.100 0.765 -0.333 -0.538 -0.150 ... ]
<i>movie</i>	→	[ 0.104 -0.054 -0.268 -0.877 0.005 ... ]
<i>!</i>	→	[ 0.439 -0.577 -0.727 0.261 0.699 ... ]

word vectors  
word embeddings  
neural embeddings  
dense embeddings  
others...

- Similarity of vectors represents similarity of meaning for particular words

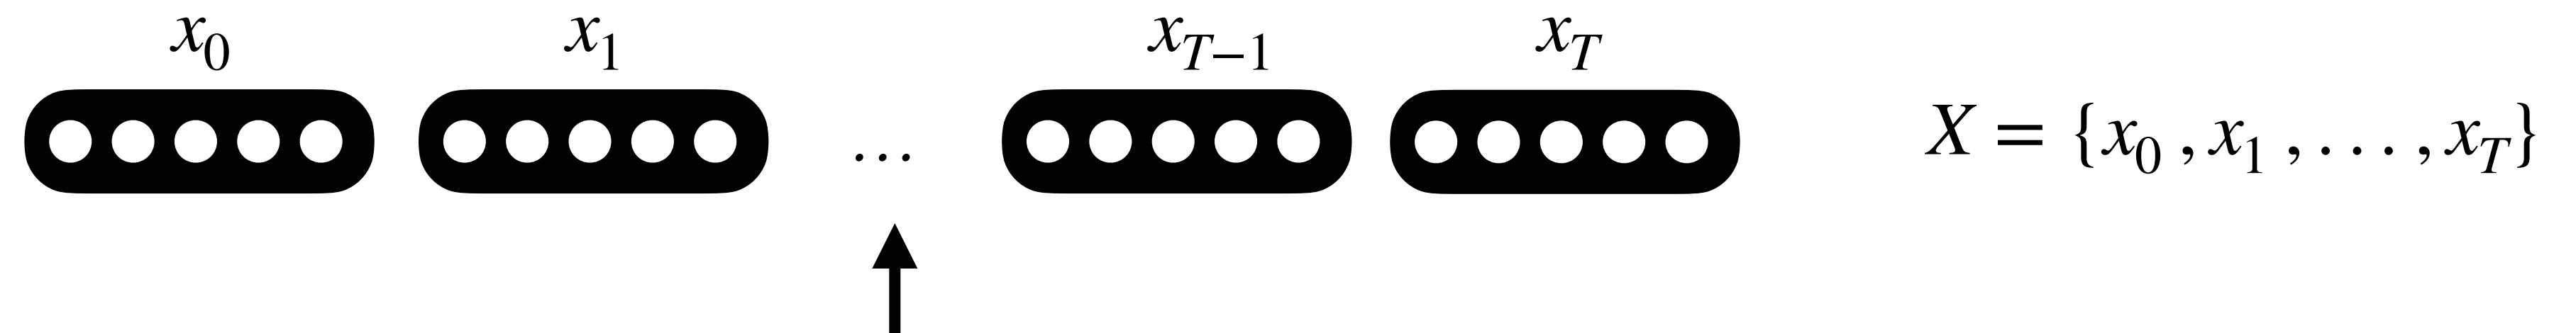
# A simple NLP model

- For each sequence  $S$ , we have a corresponding sequence of embeddings  $X$



# A simple NLP model

- For each sequence  $S$ , we have a corresponding sequence of embeddings  $X$



$S_1 =$  *I **really** enjoyed the movie **we** watched on Saturday !*

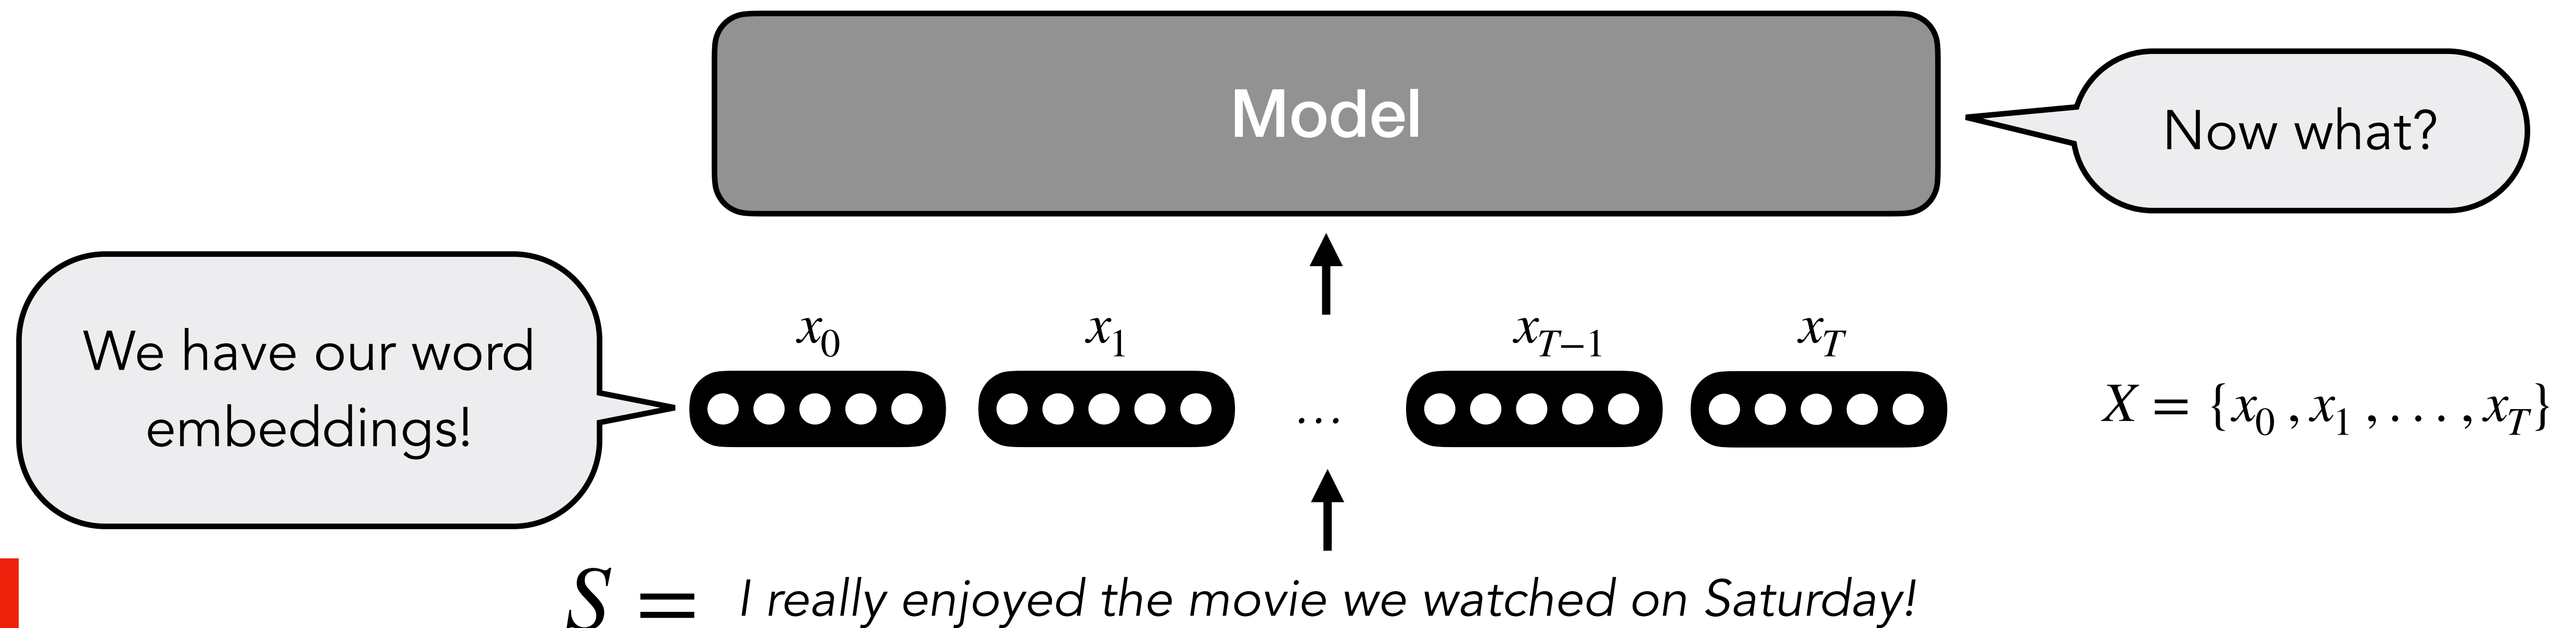
- Embeddings  $x_t \in X$  are indexed from shared embedding dictionary  $\mathbb{E}$  for all items in vocabulary  $V$

$S_2 =$  *We **really** loved a film **we** saw last Sunday !*

Bolded words would index the same embedding in  $\mathbb{E}$

# A simple NLP model

- For each sequence  $S$ , we have a corresponding sequence of embeddings  $X$



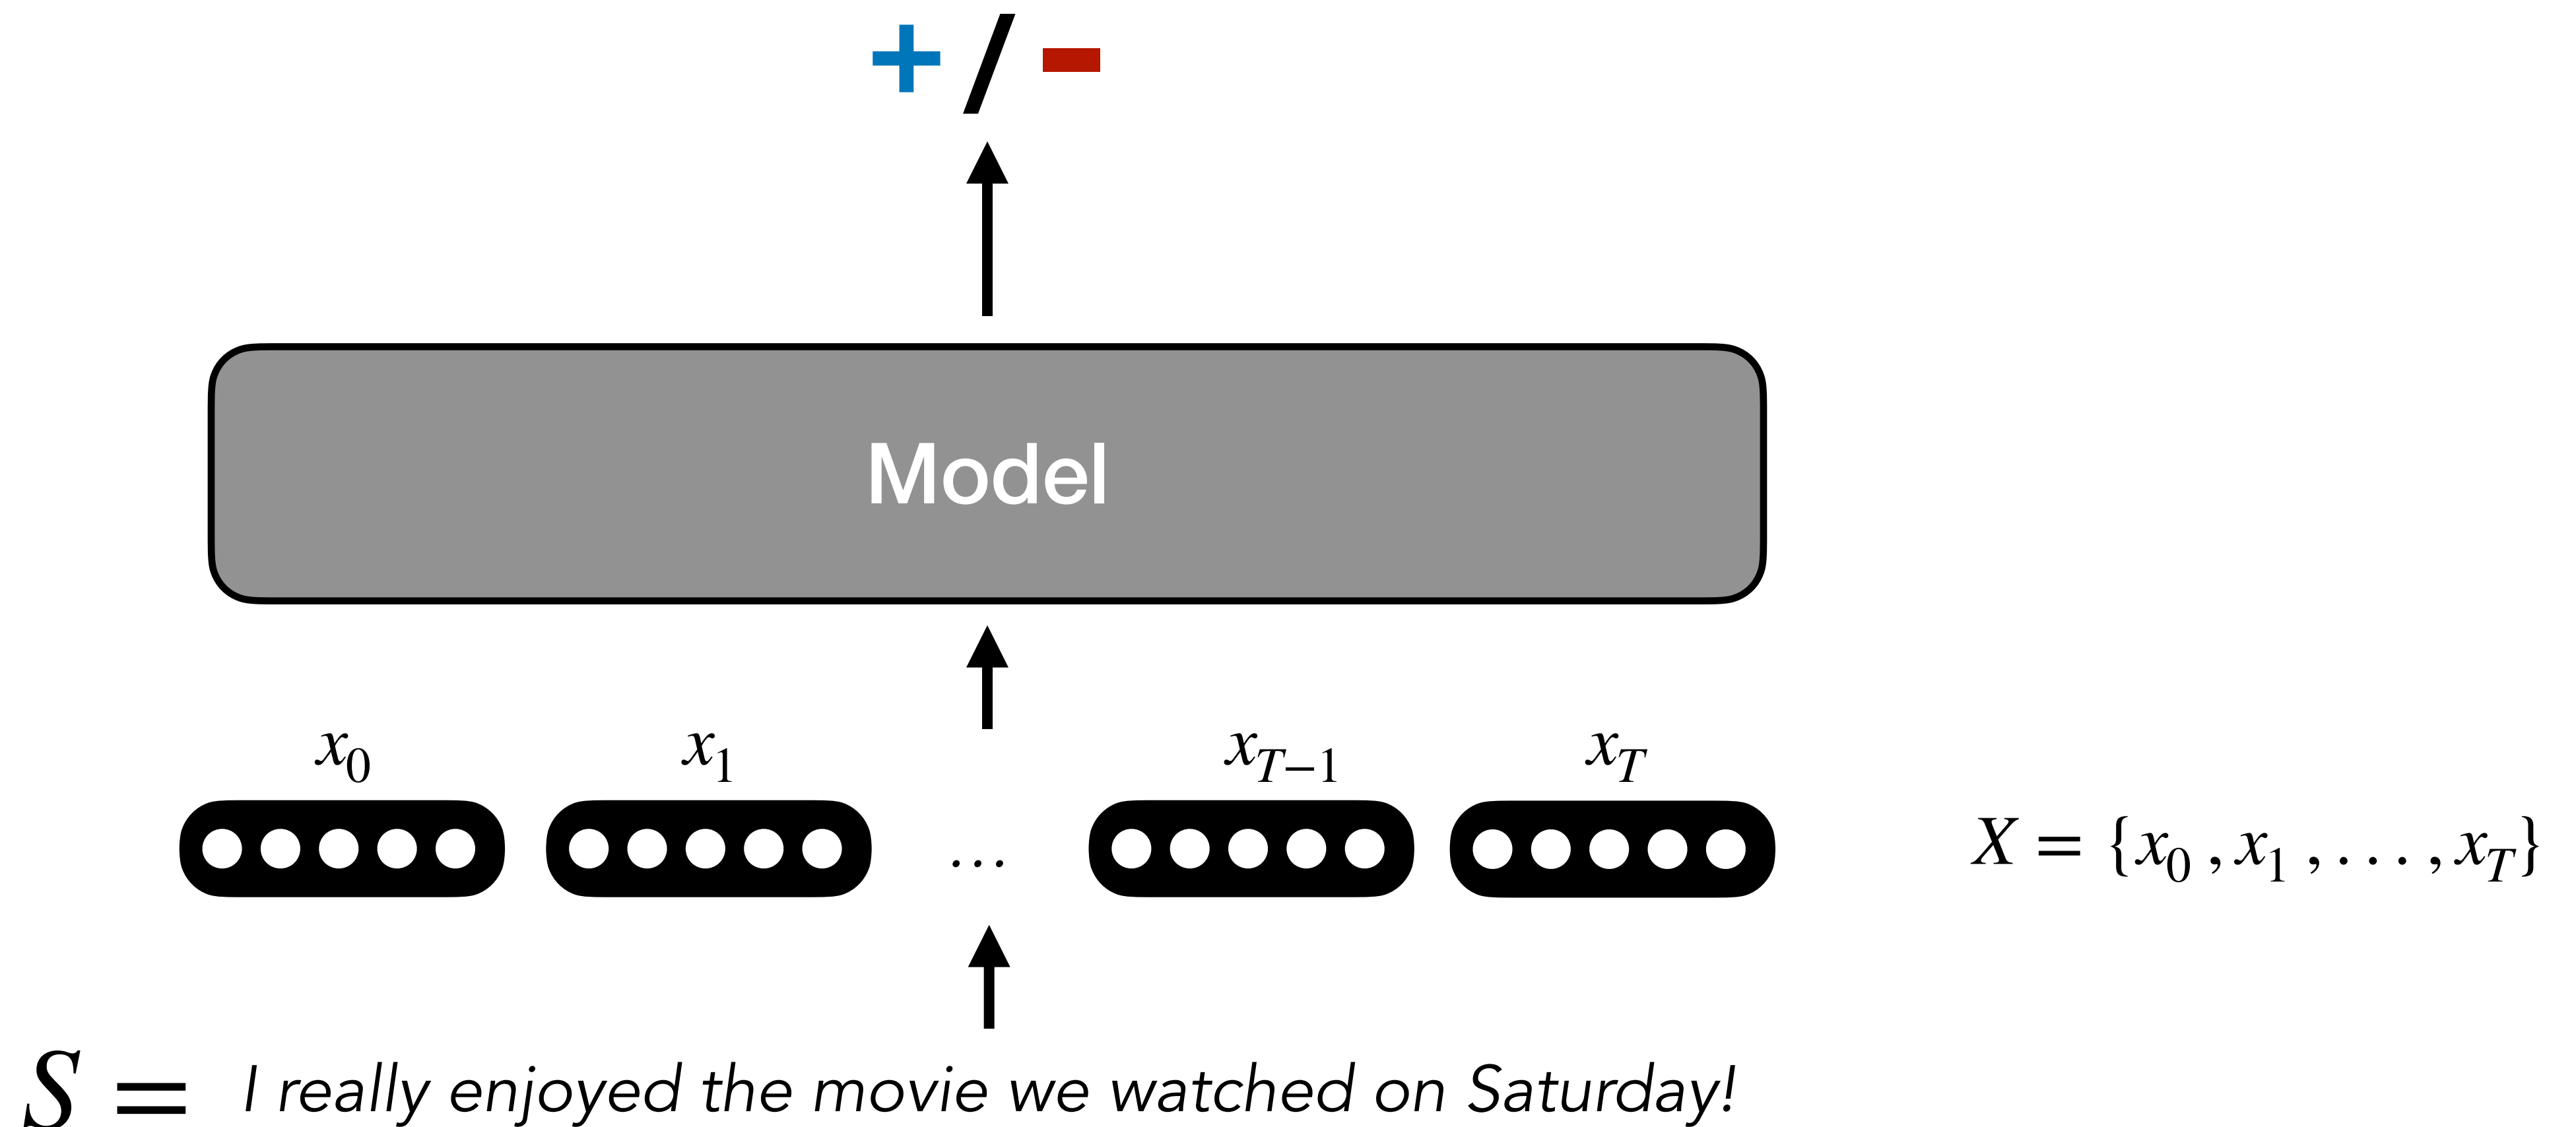


# Question

**What should we use as a model?**

# A simple NLP model

- Our model modifies and / or composes these word embeddings to formulate a representation that allows it to predict the correct label



# A simple NLP model

- Our model modifies and / or composes these word embeddings to formulate a representation that allows it to predict the correct label
  - Recurrent neural networks (RNNs) and variants (LSTM, GRU) - Week 2
  - Self-attention & Transformer - Week 3
  - State-space Models (not covered in this course)
  - Multiple of the above ?

# A simple NLP model


- Our model modifies and / or composes these word embeddings to formulate a representation that allows it to predict the correct label
  - Recurrent neural networks (RNNs) and variants (LSTM, GRU) - Week 2
  - Self-attention & Transformer - Week 3
  - State-space Models (not covered in this course)
  - Multiple of the above ?
  - Or perhaps something super simple: **Sum-pool, Avg-pool, Max-pool?**

# A simple NLP model

Notation: Typically, we represent the output of a model as  $h$  (or  $o$ ).

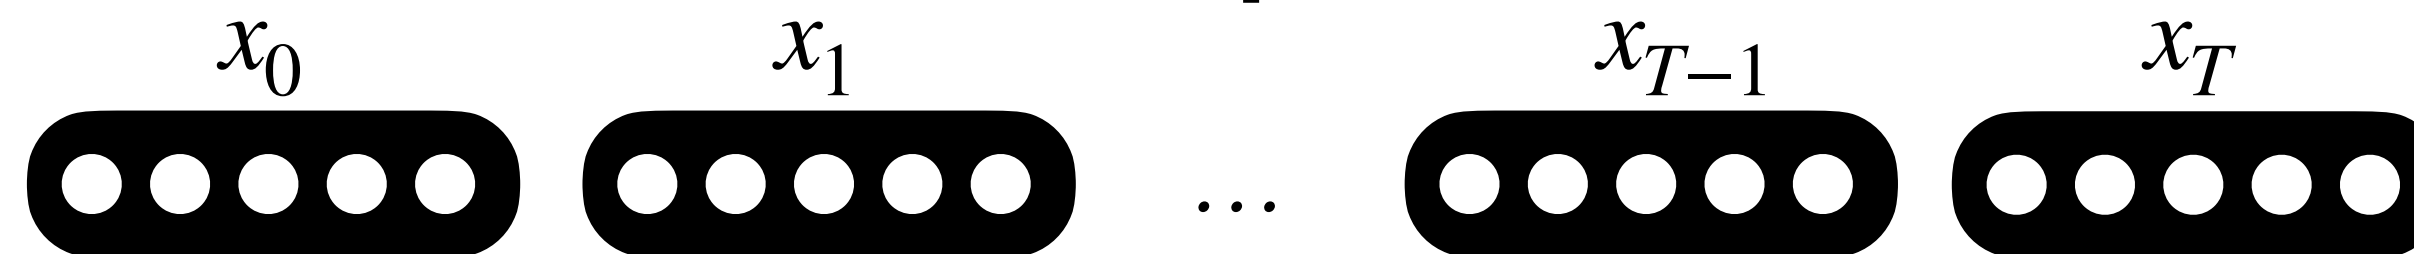
$$h_T = \sum_{t=0}^T x_t$$

+ / -



We composed our embeddings into a different representation!

Sum-pool



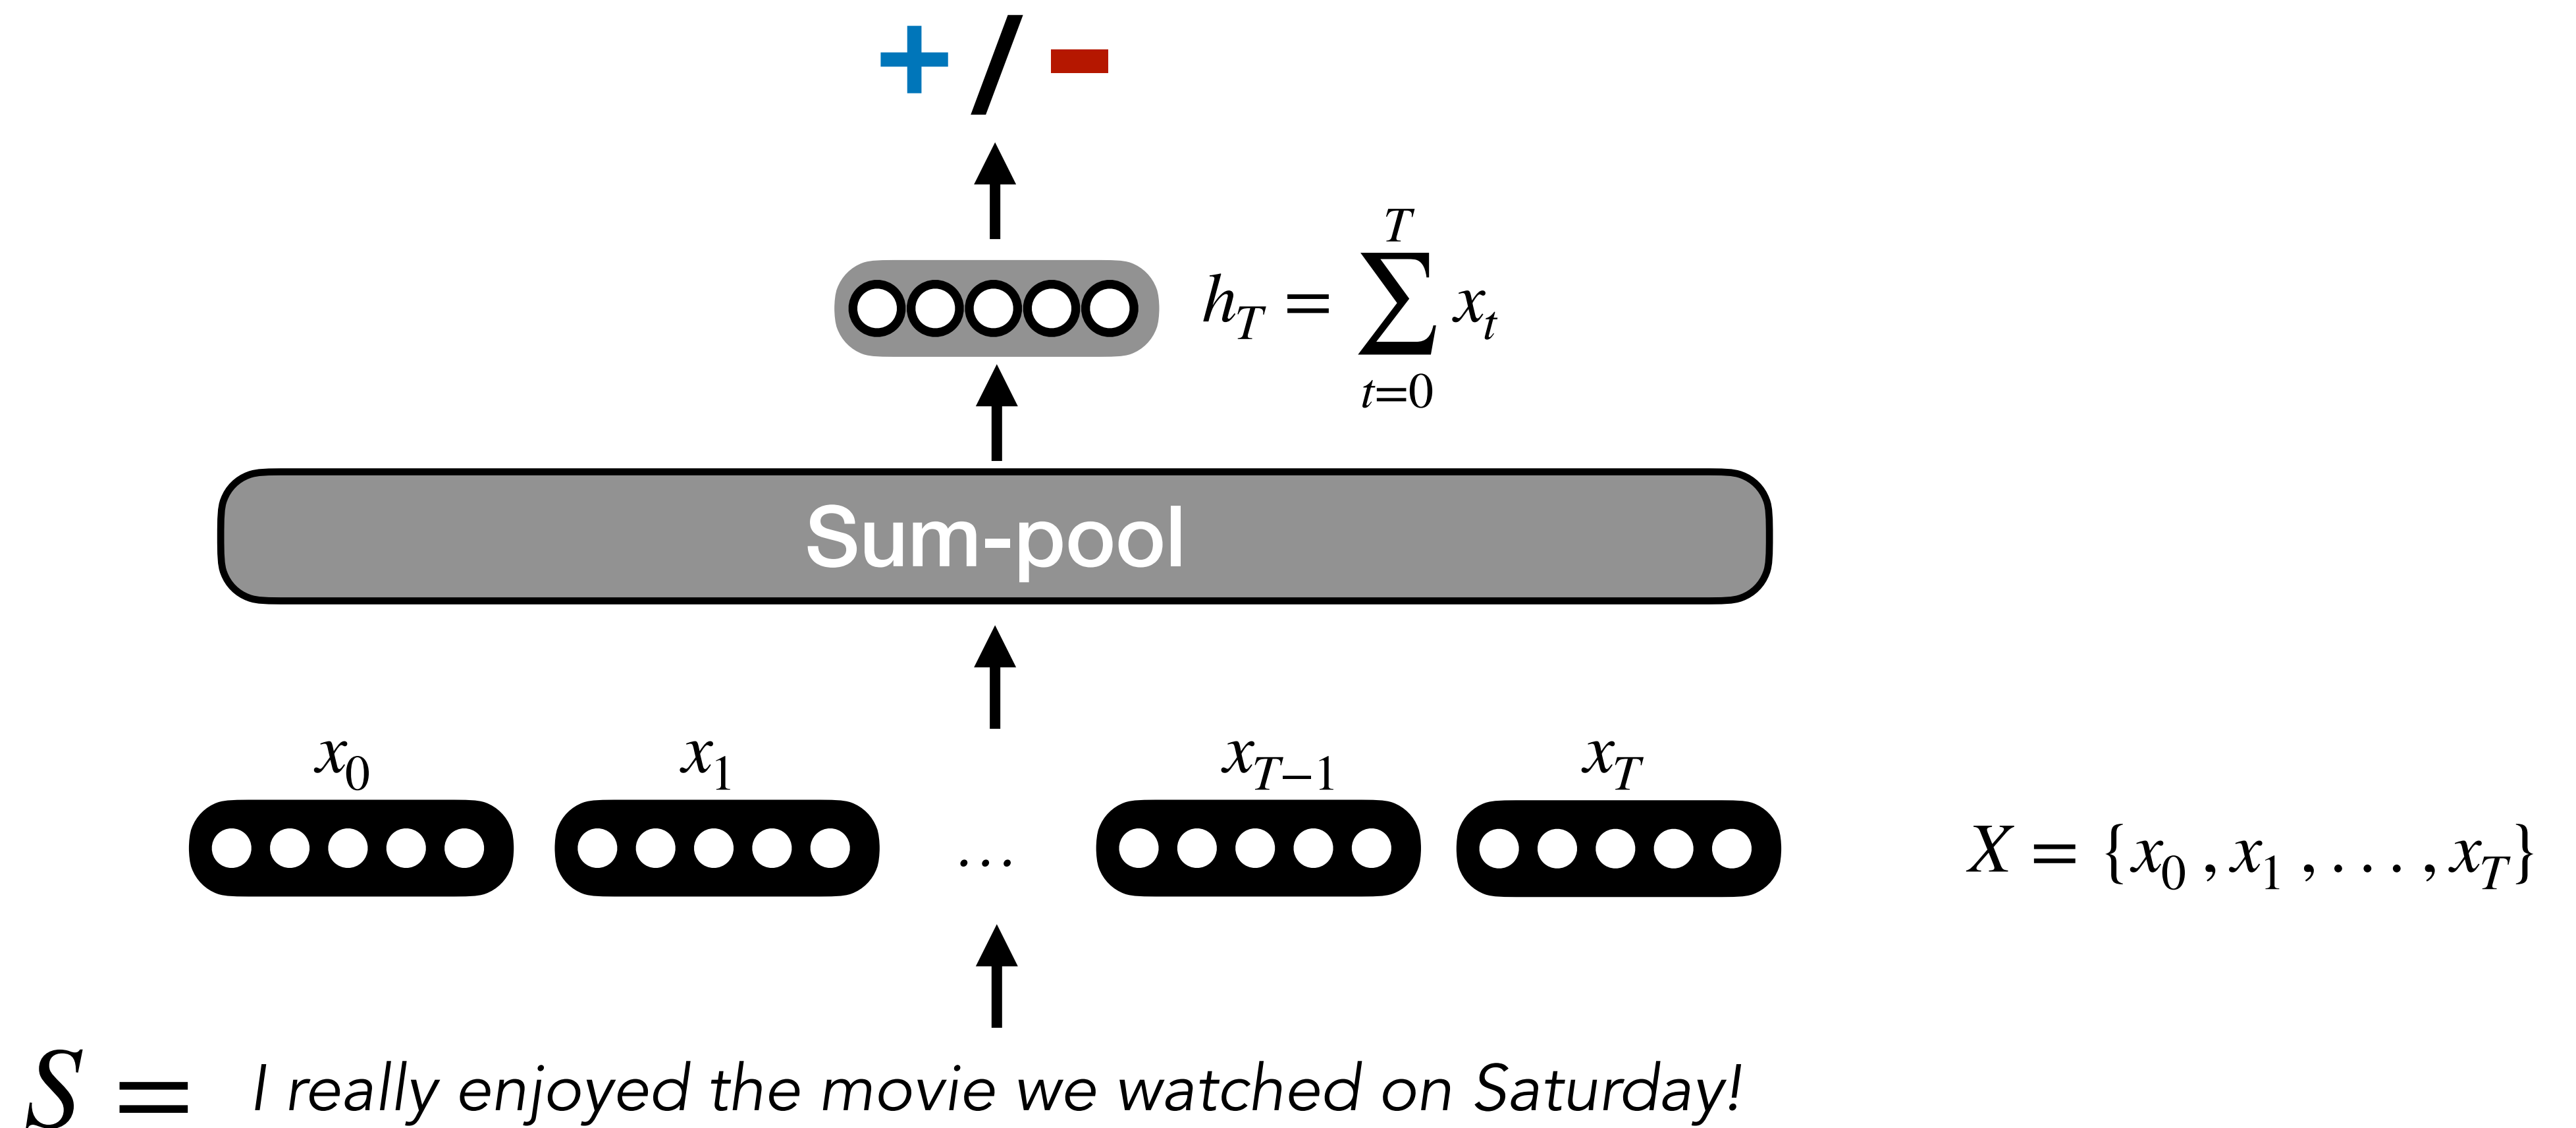
$$X = \{x_0, x_1, \dots, x_T\}$$

$S =$  *I really enjoyed the movie we watched on Saturday!*

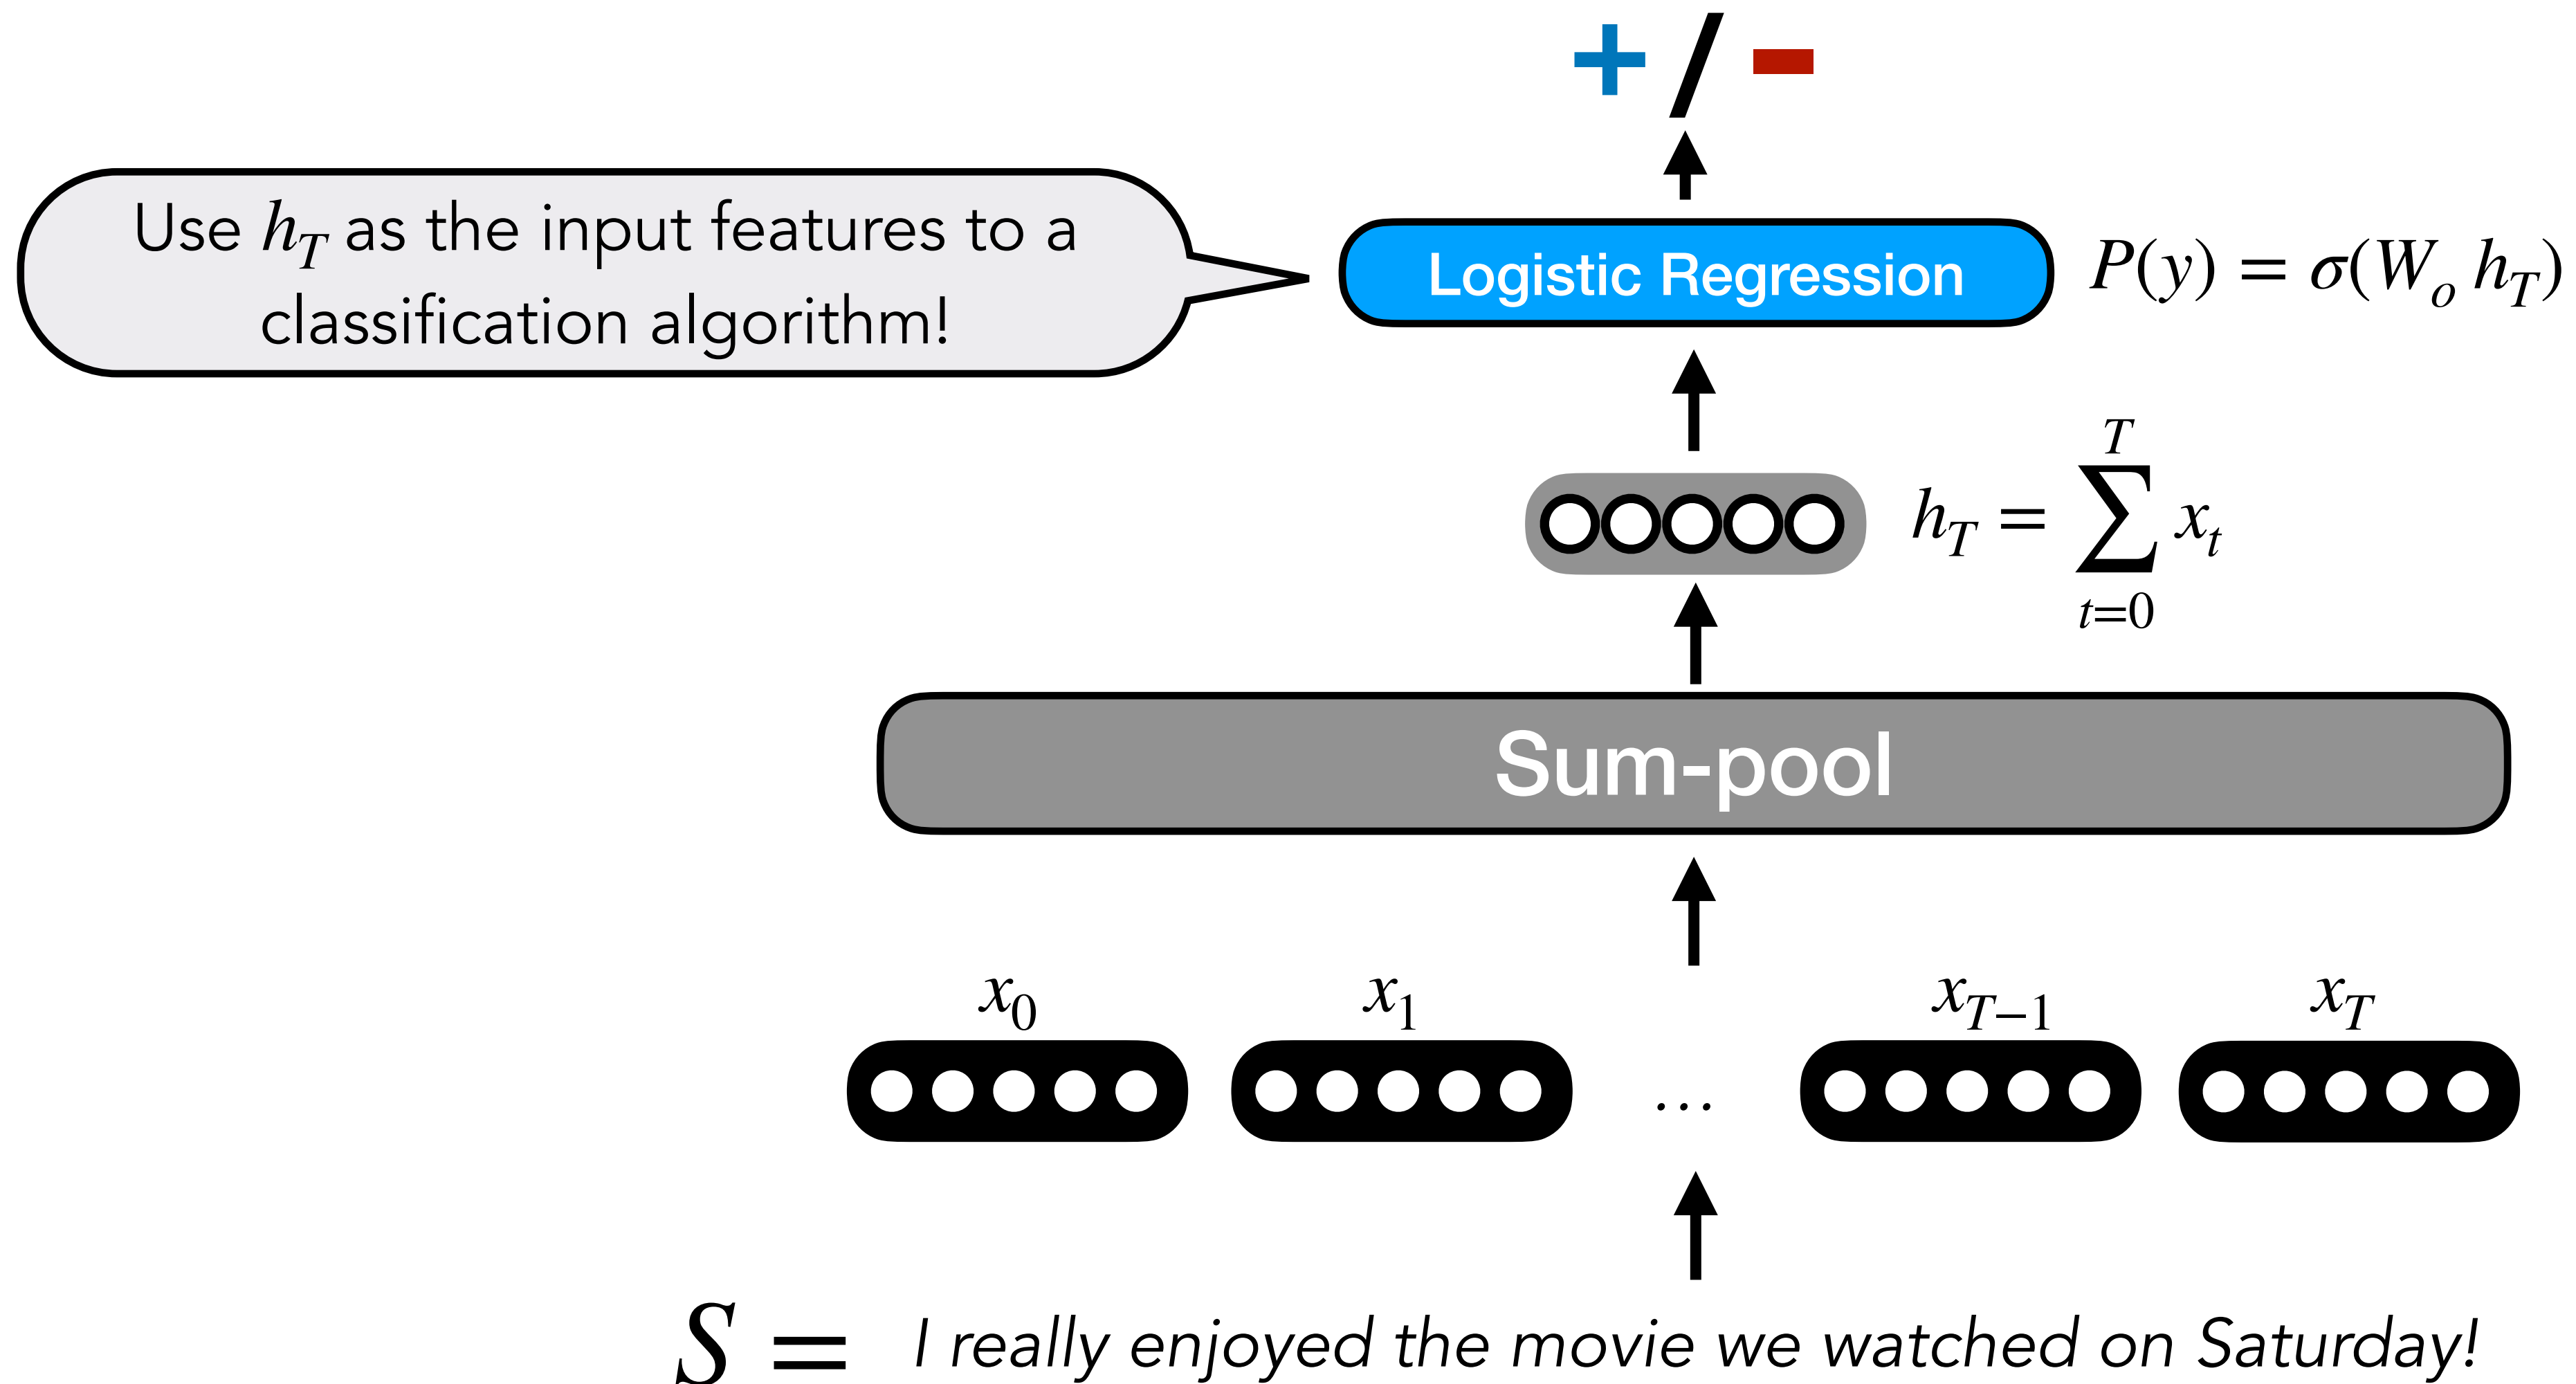
# Question

**How do we convert the output of our model to a prediction?**

# Predicting the label



# Predicting the label



Learn using  
**backpropagation:**  
compute gradients of  
loss with respect to  
initial embeddings  $X$

Learn embeddings  
that allow you to do  
the task successfully!

$$X = \{x_0, x_1, \dots, x_T\}$$




# Question

**How could we use our model for tasks beyond classification?**

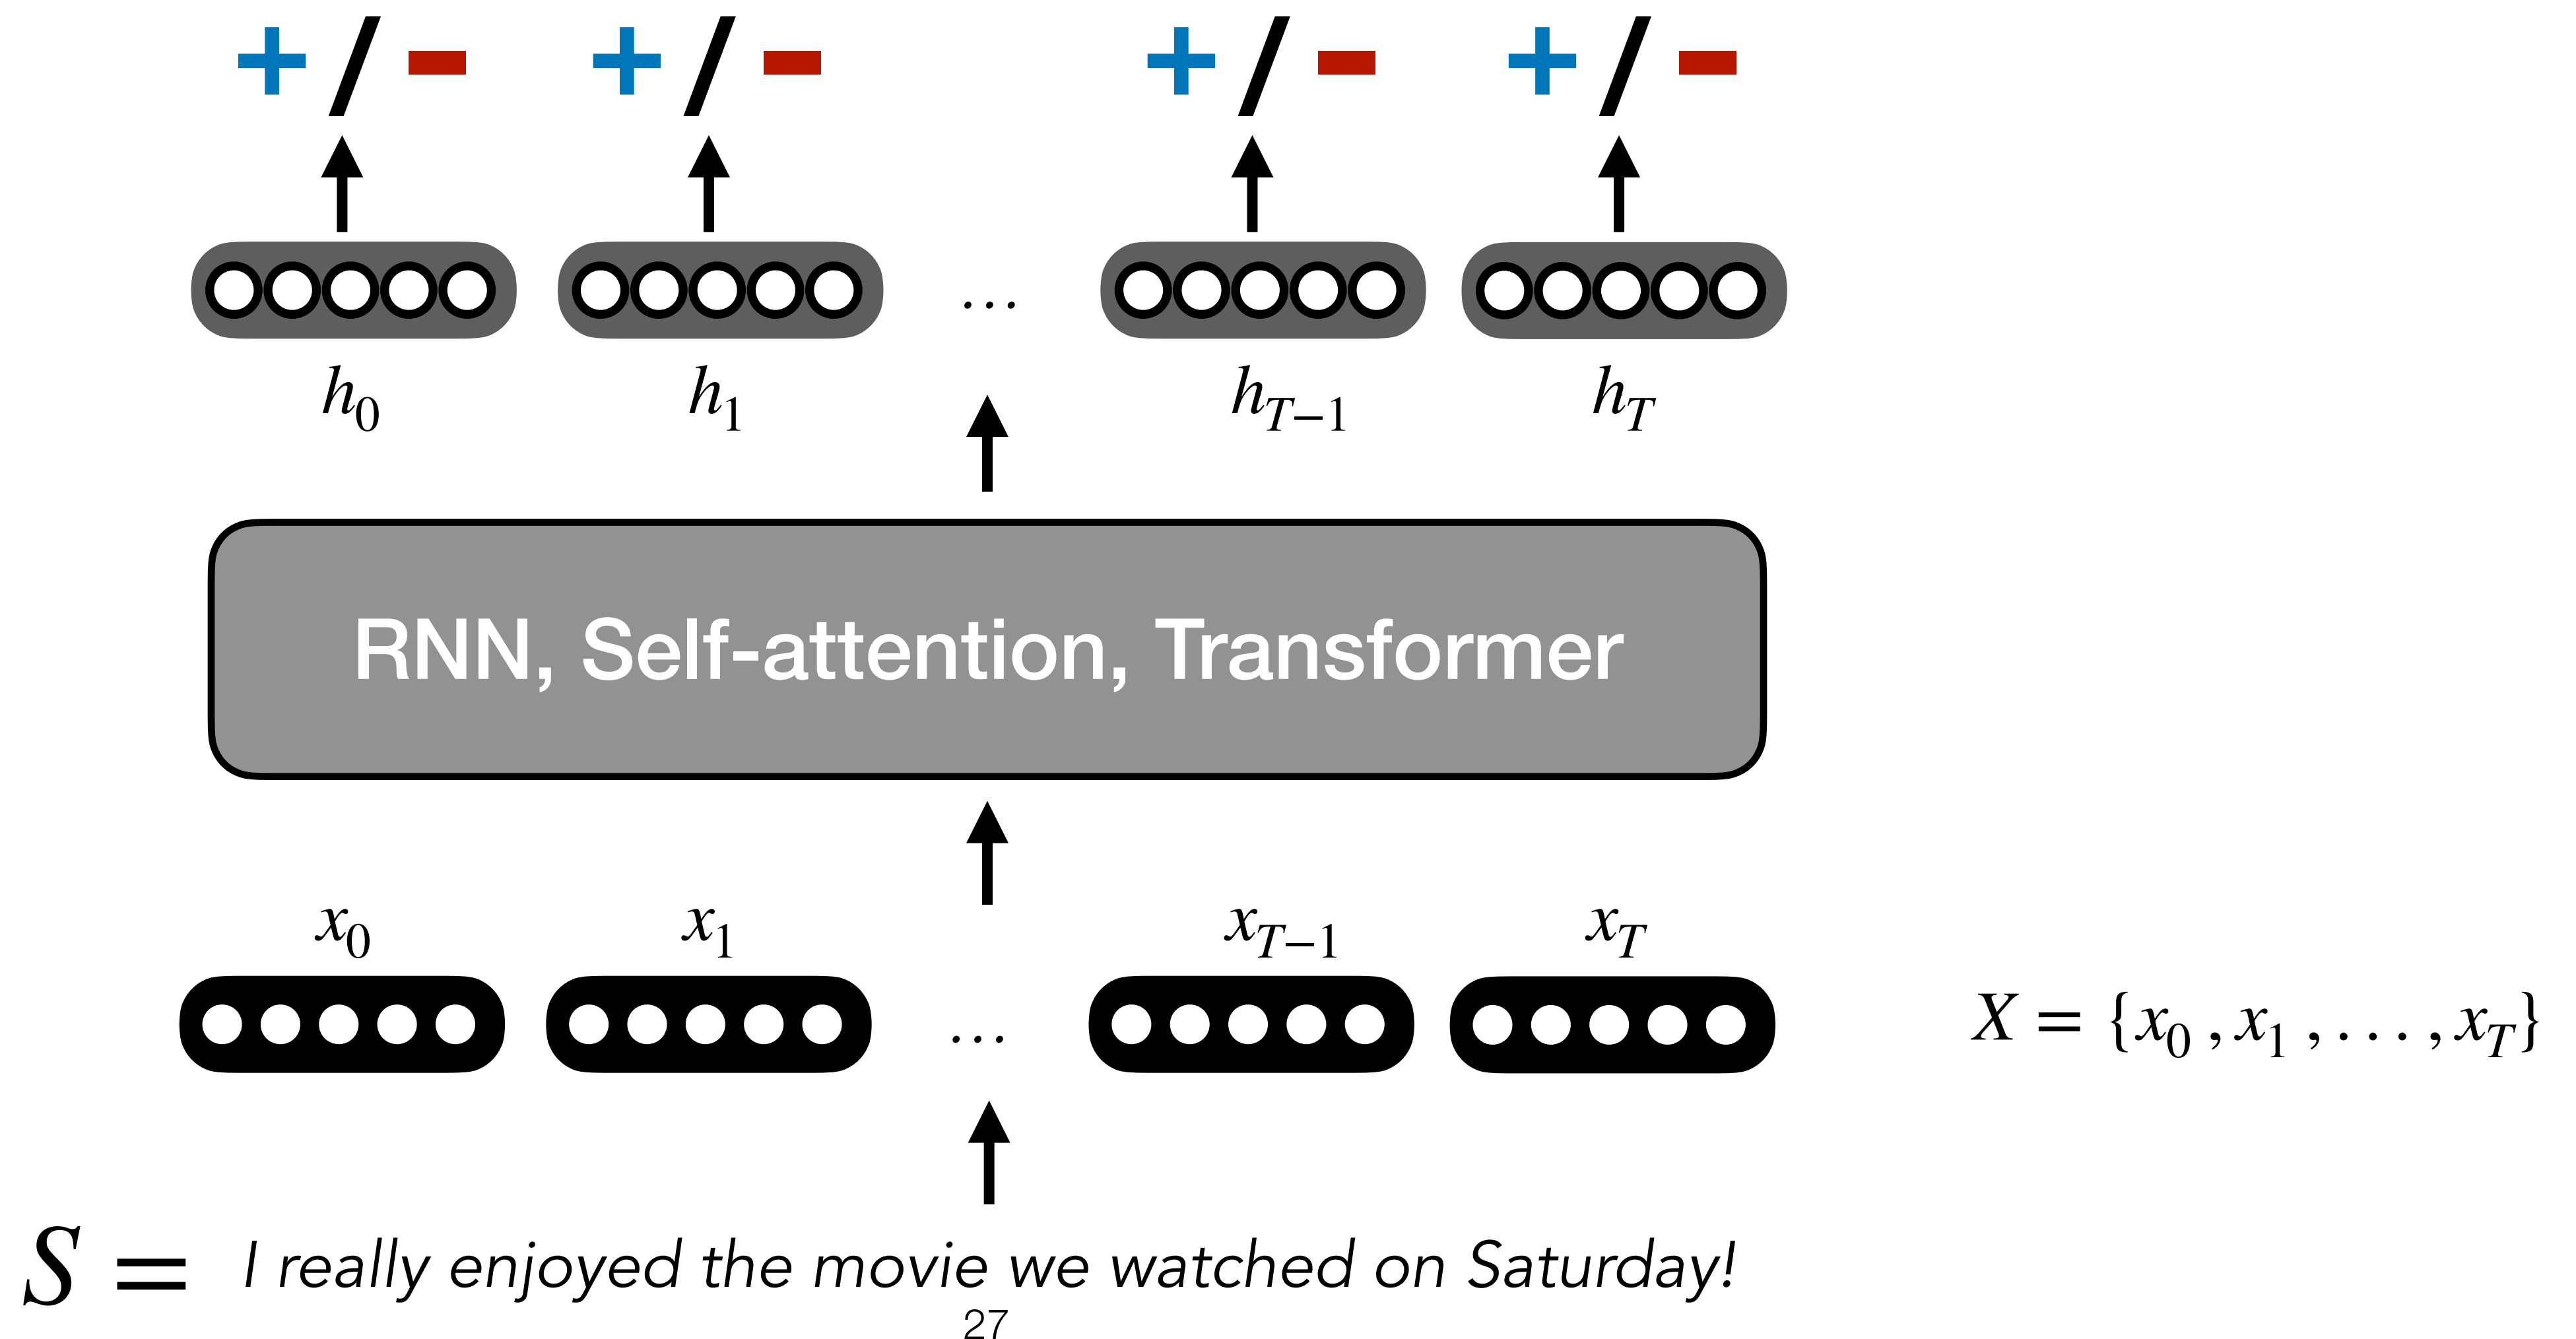
# Sequence Labeling

- **Example:** Identify which words correspond to sentimental words



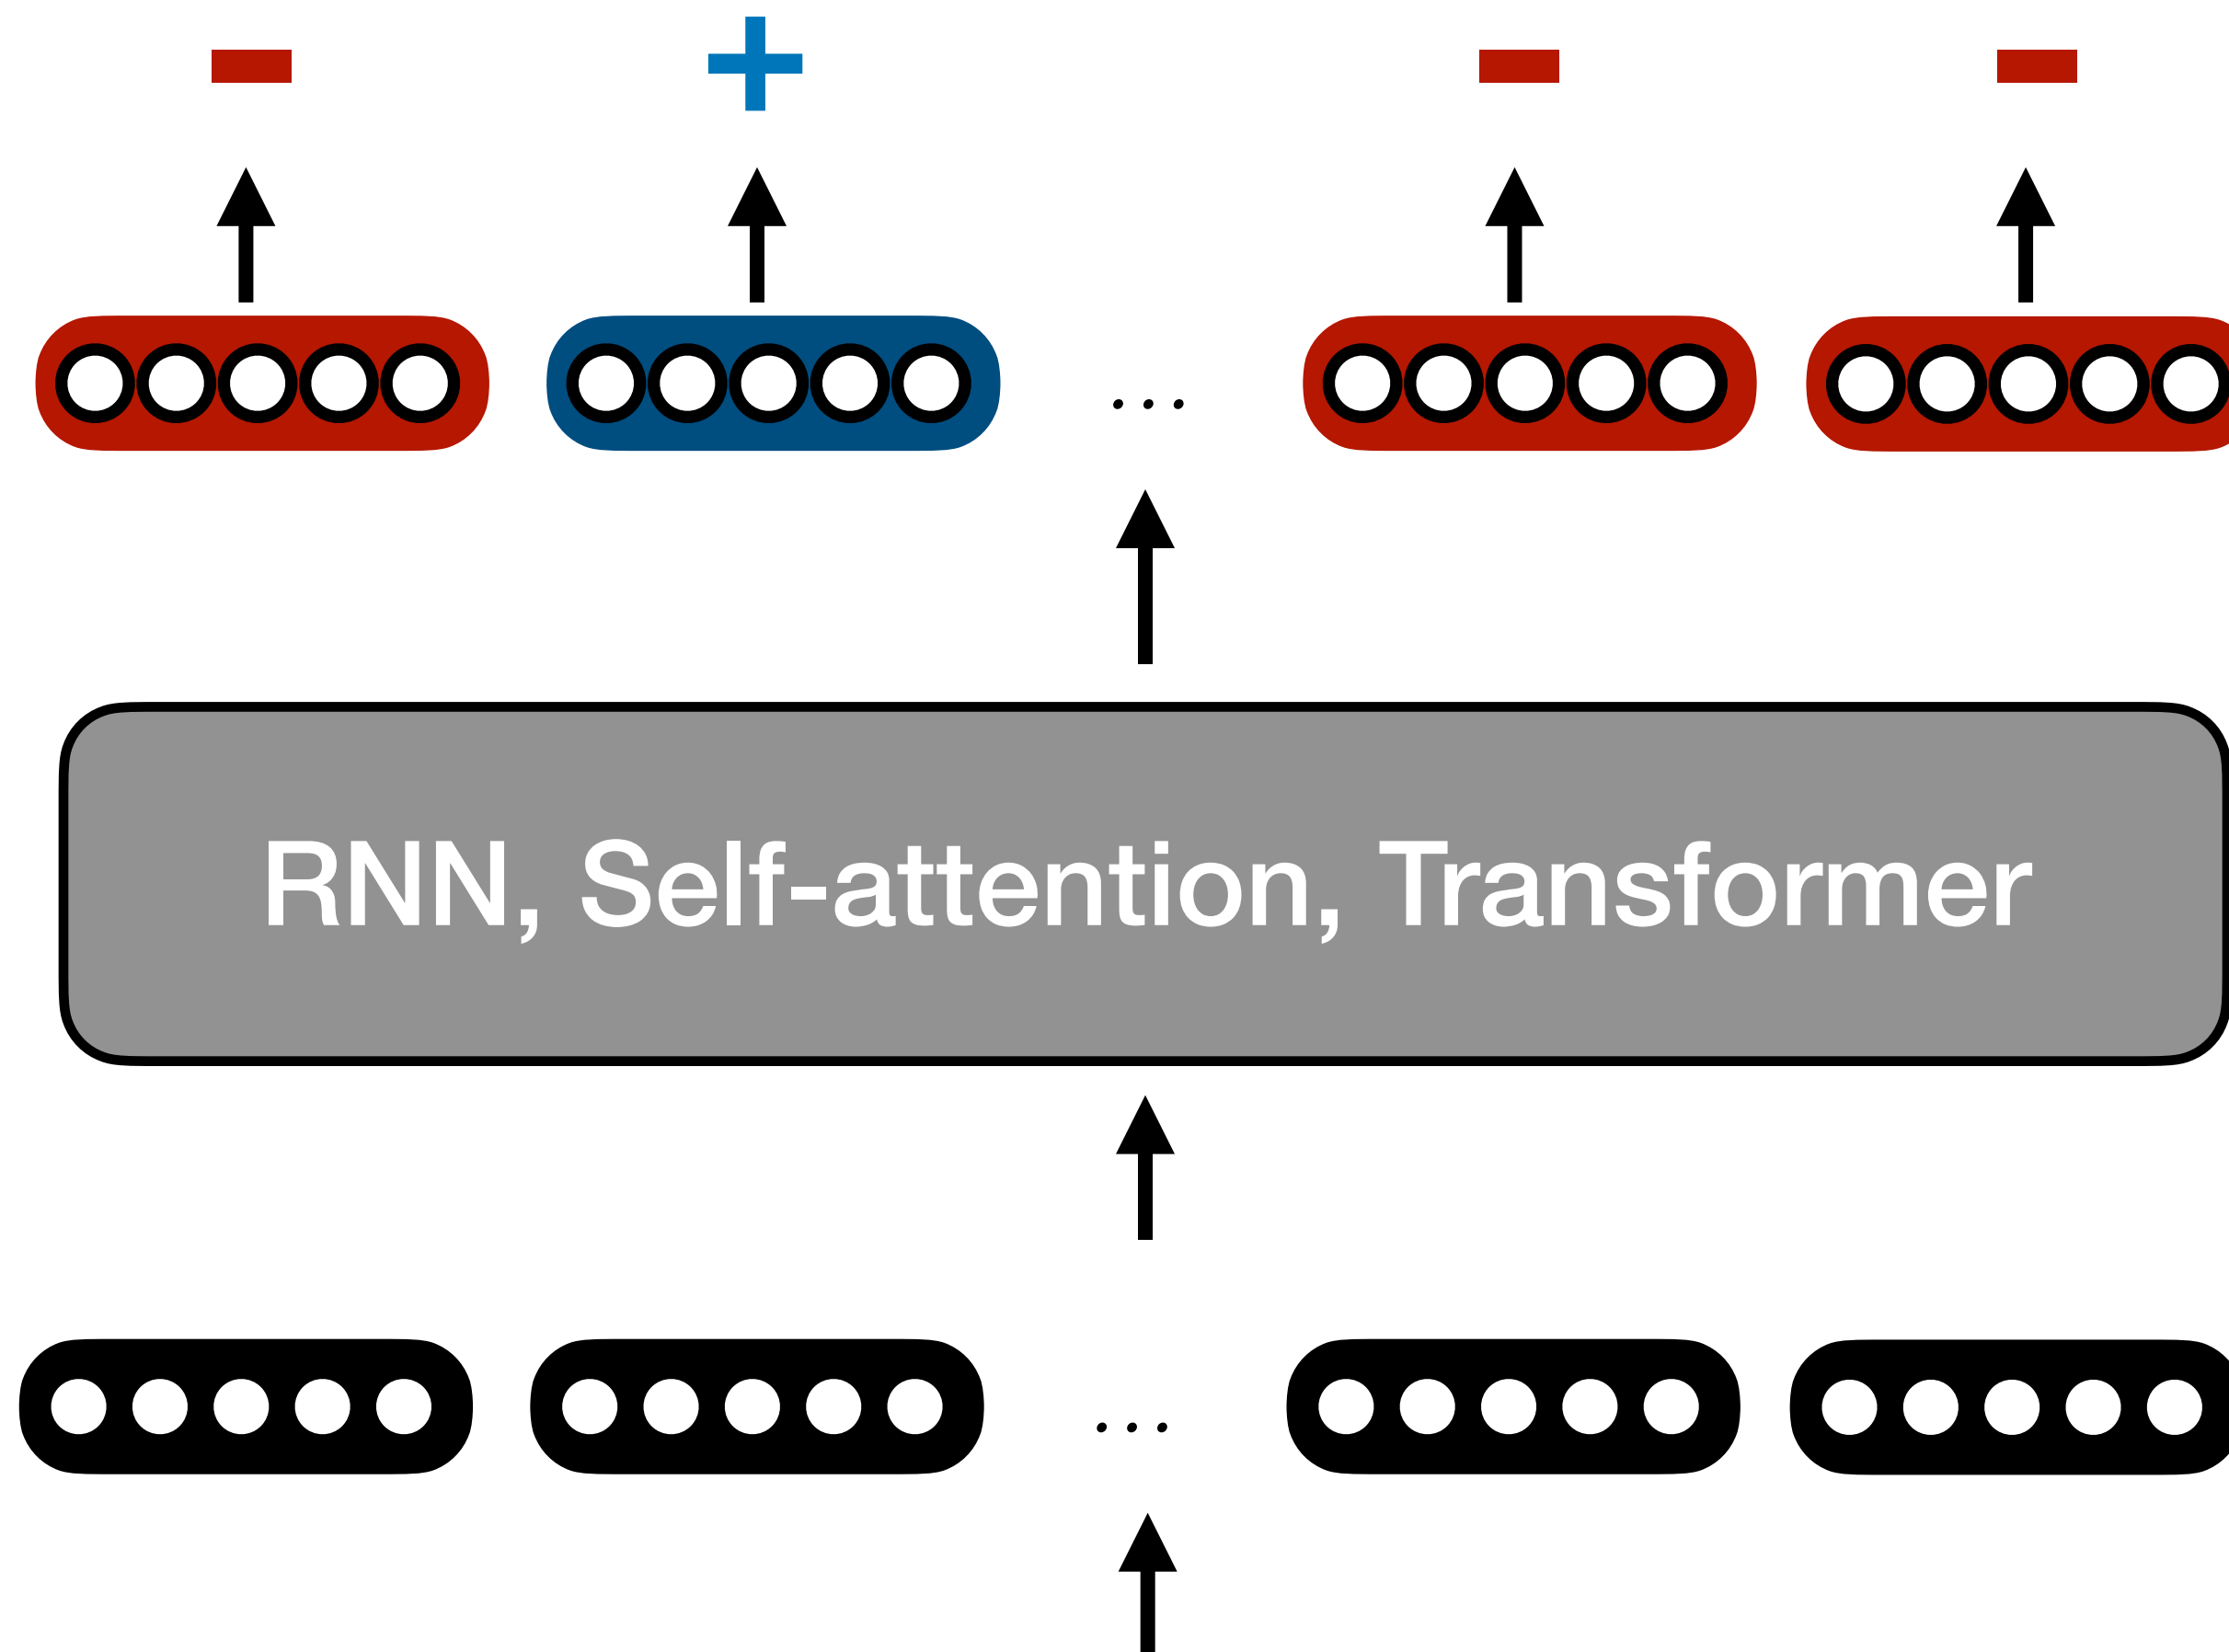
# Sequence Labeling

- **Example:** Identify which words correspond to sentimental words



# Sequence Labeling

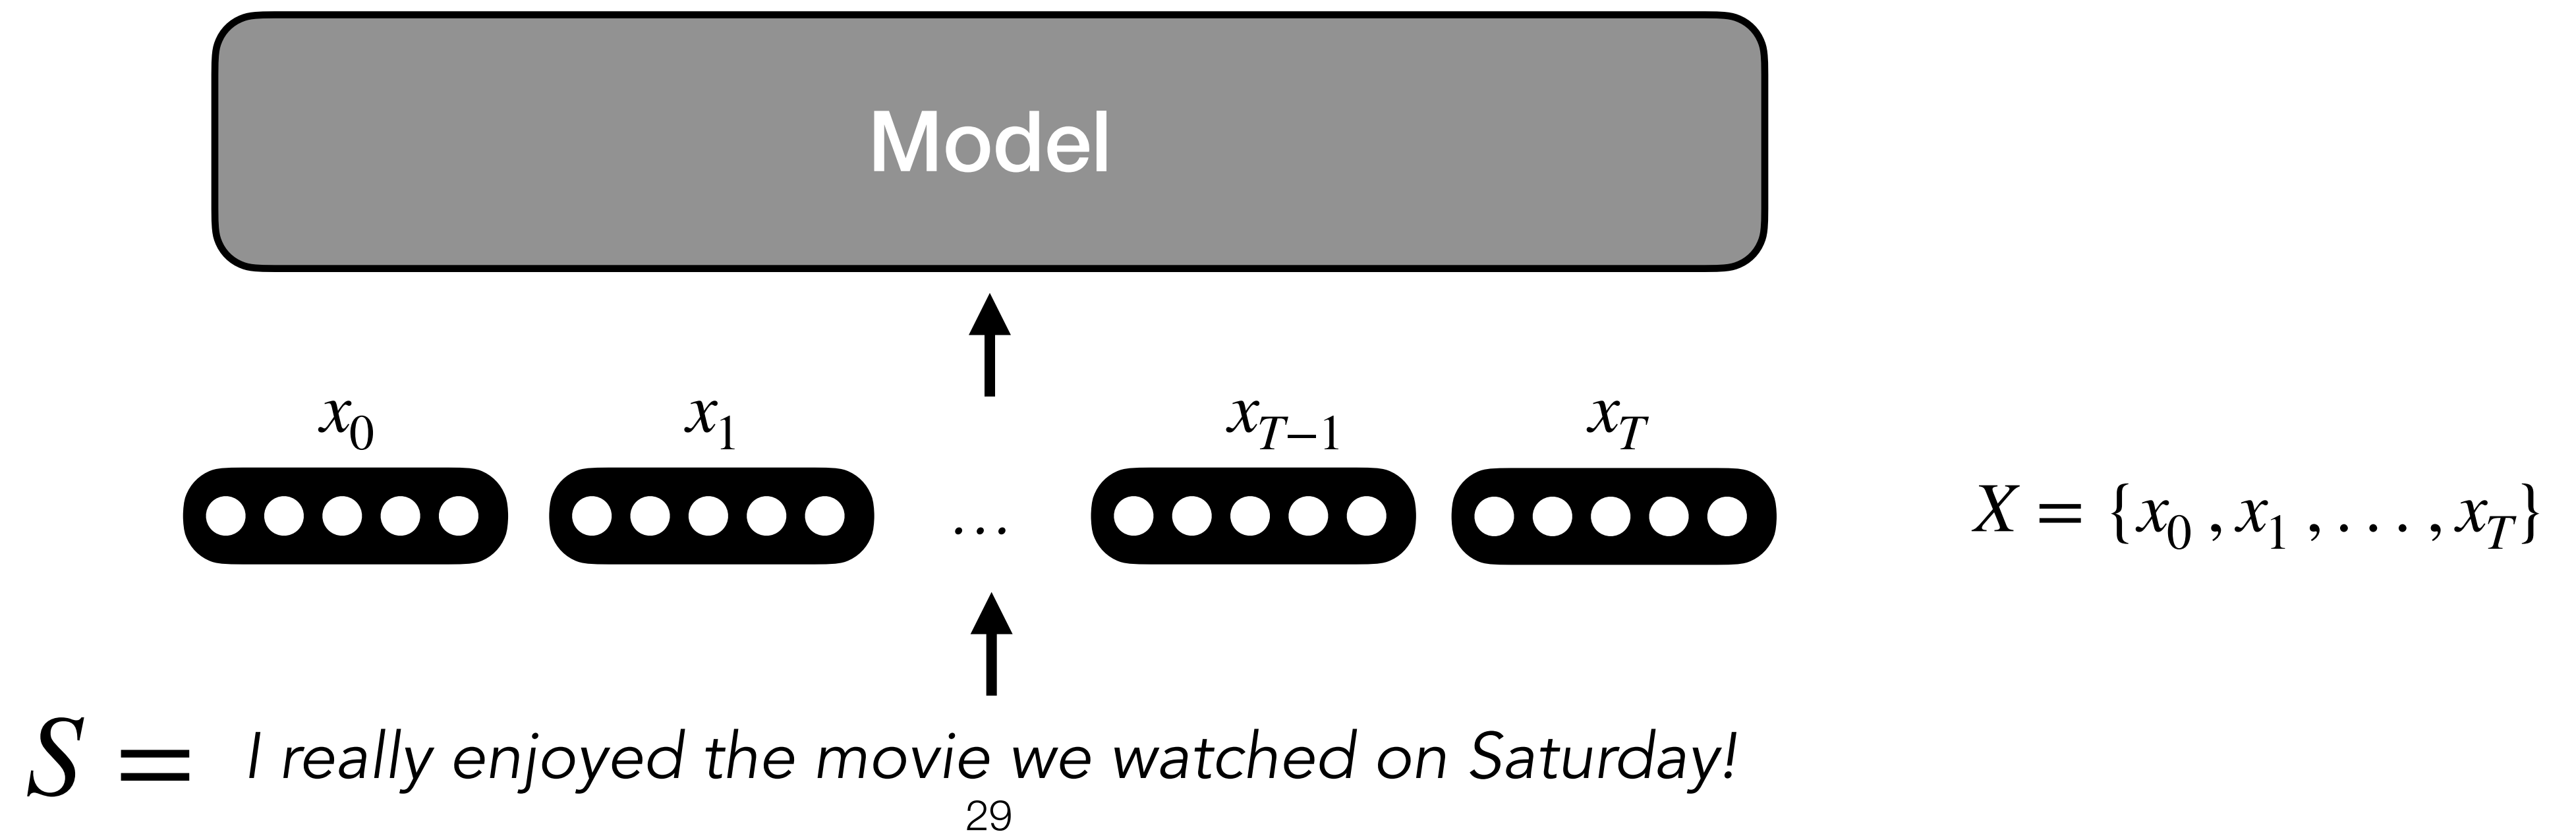
- **Example:** Identify which words correspond to sentimental words



$S =$  I really enjoyed the movie we watched on Saturday!

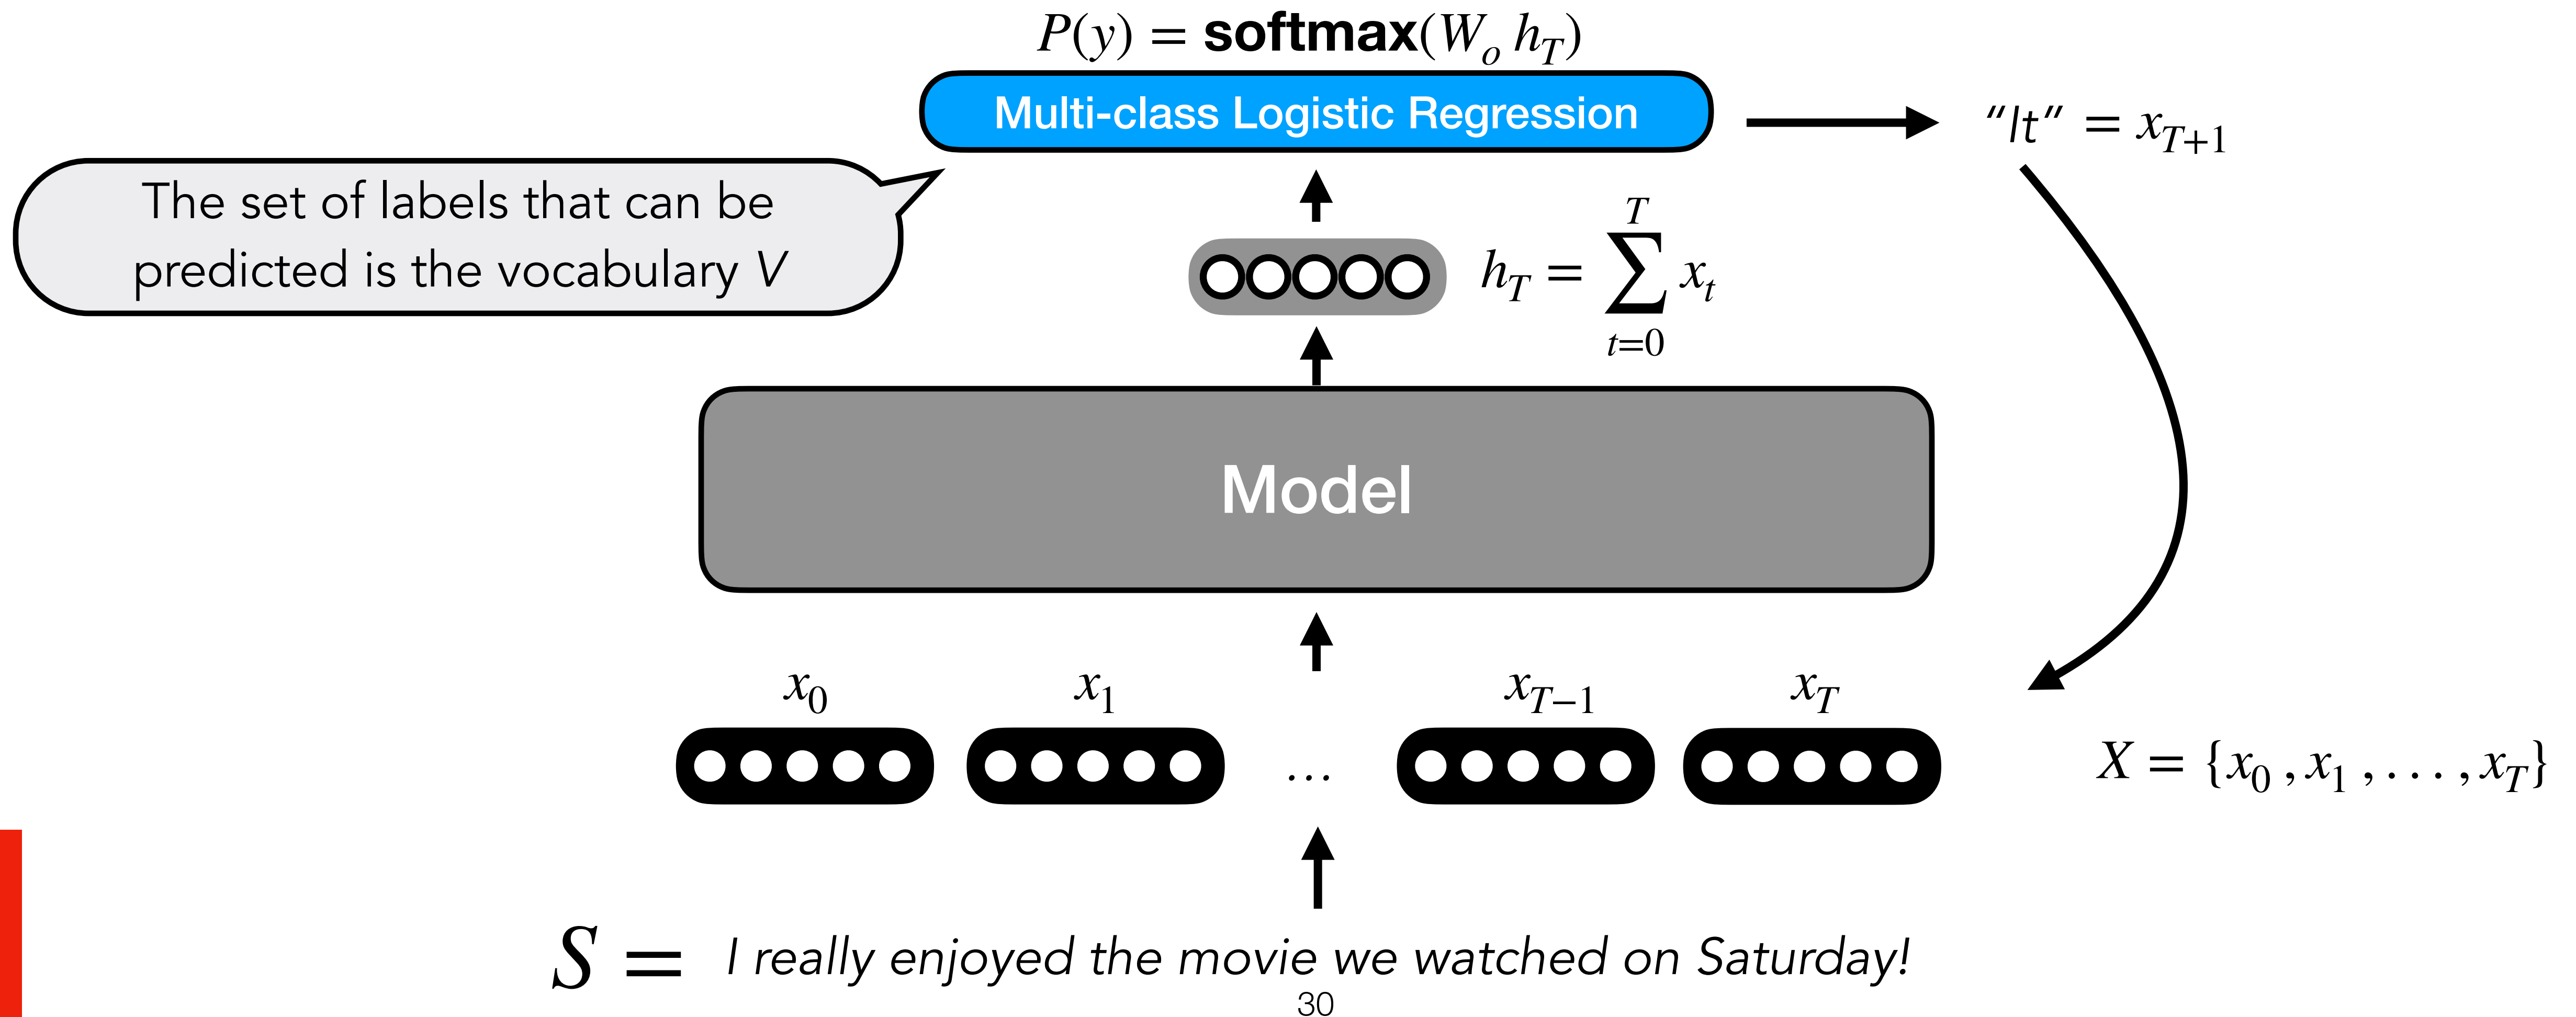
# Text Generation

- **Example:** Generate the next sentence in the review.



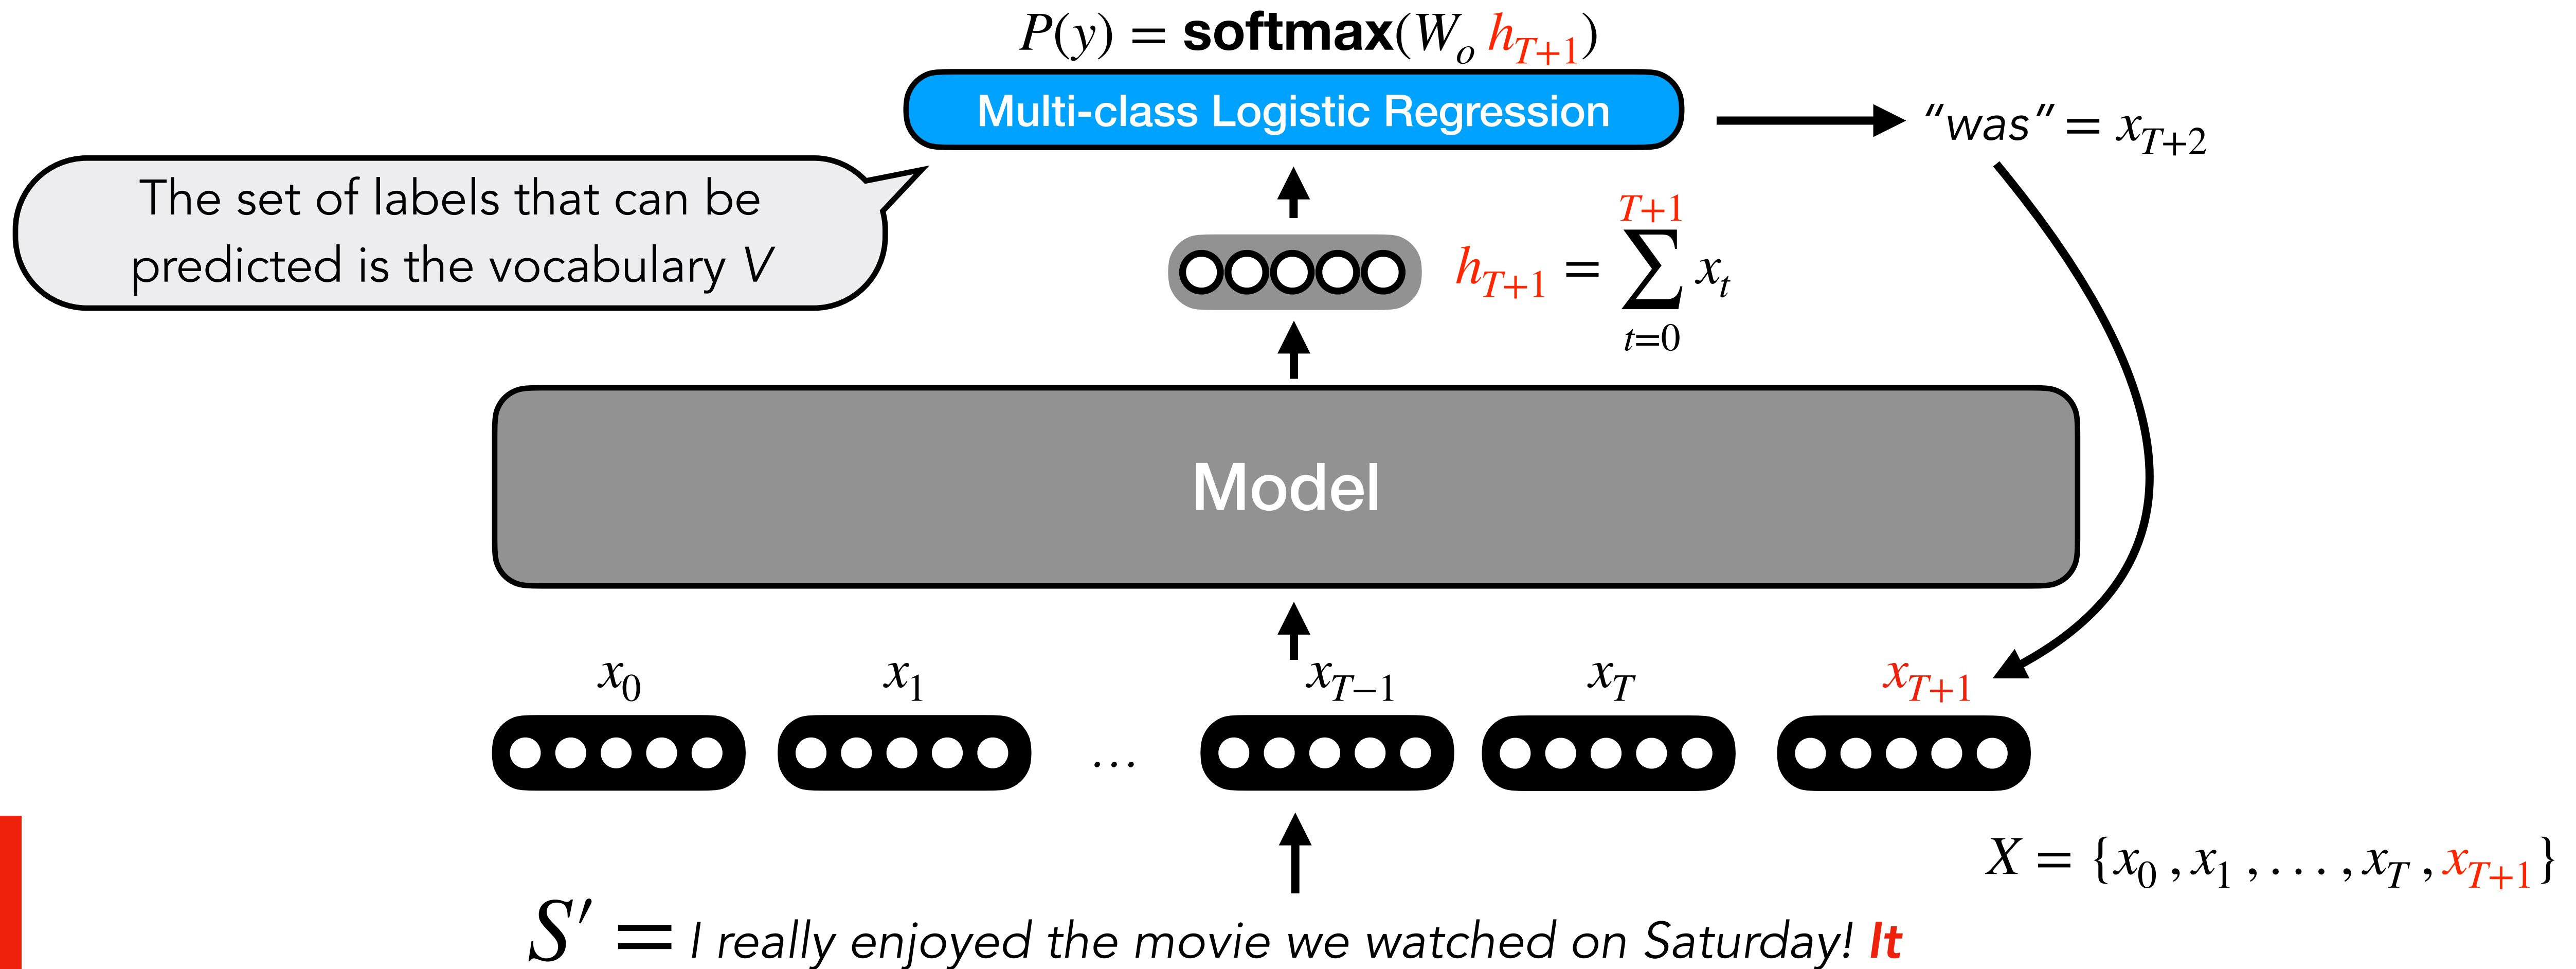
# Text Generation

- **Example:** Generate the next sentence in the review. **Word-by-word!**



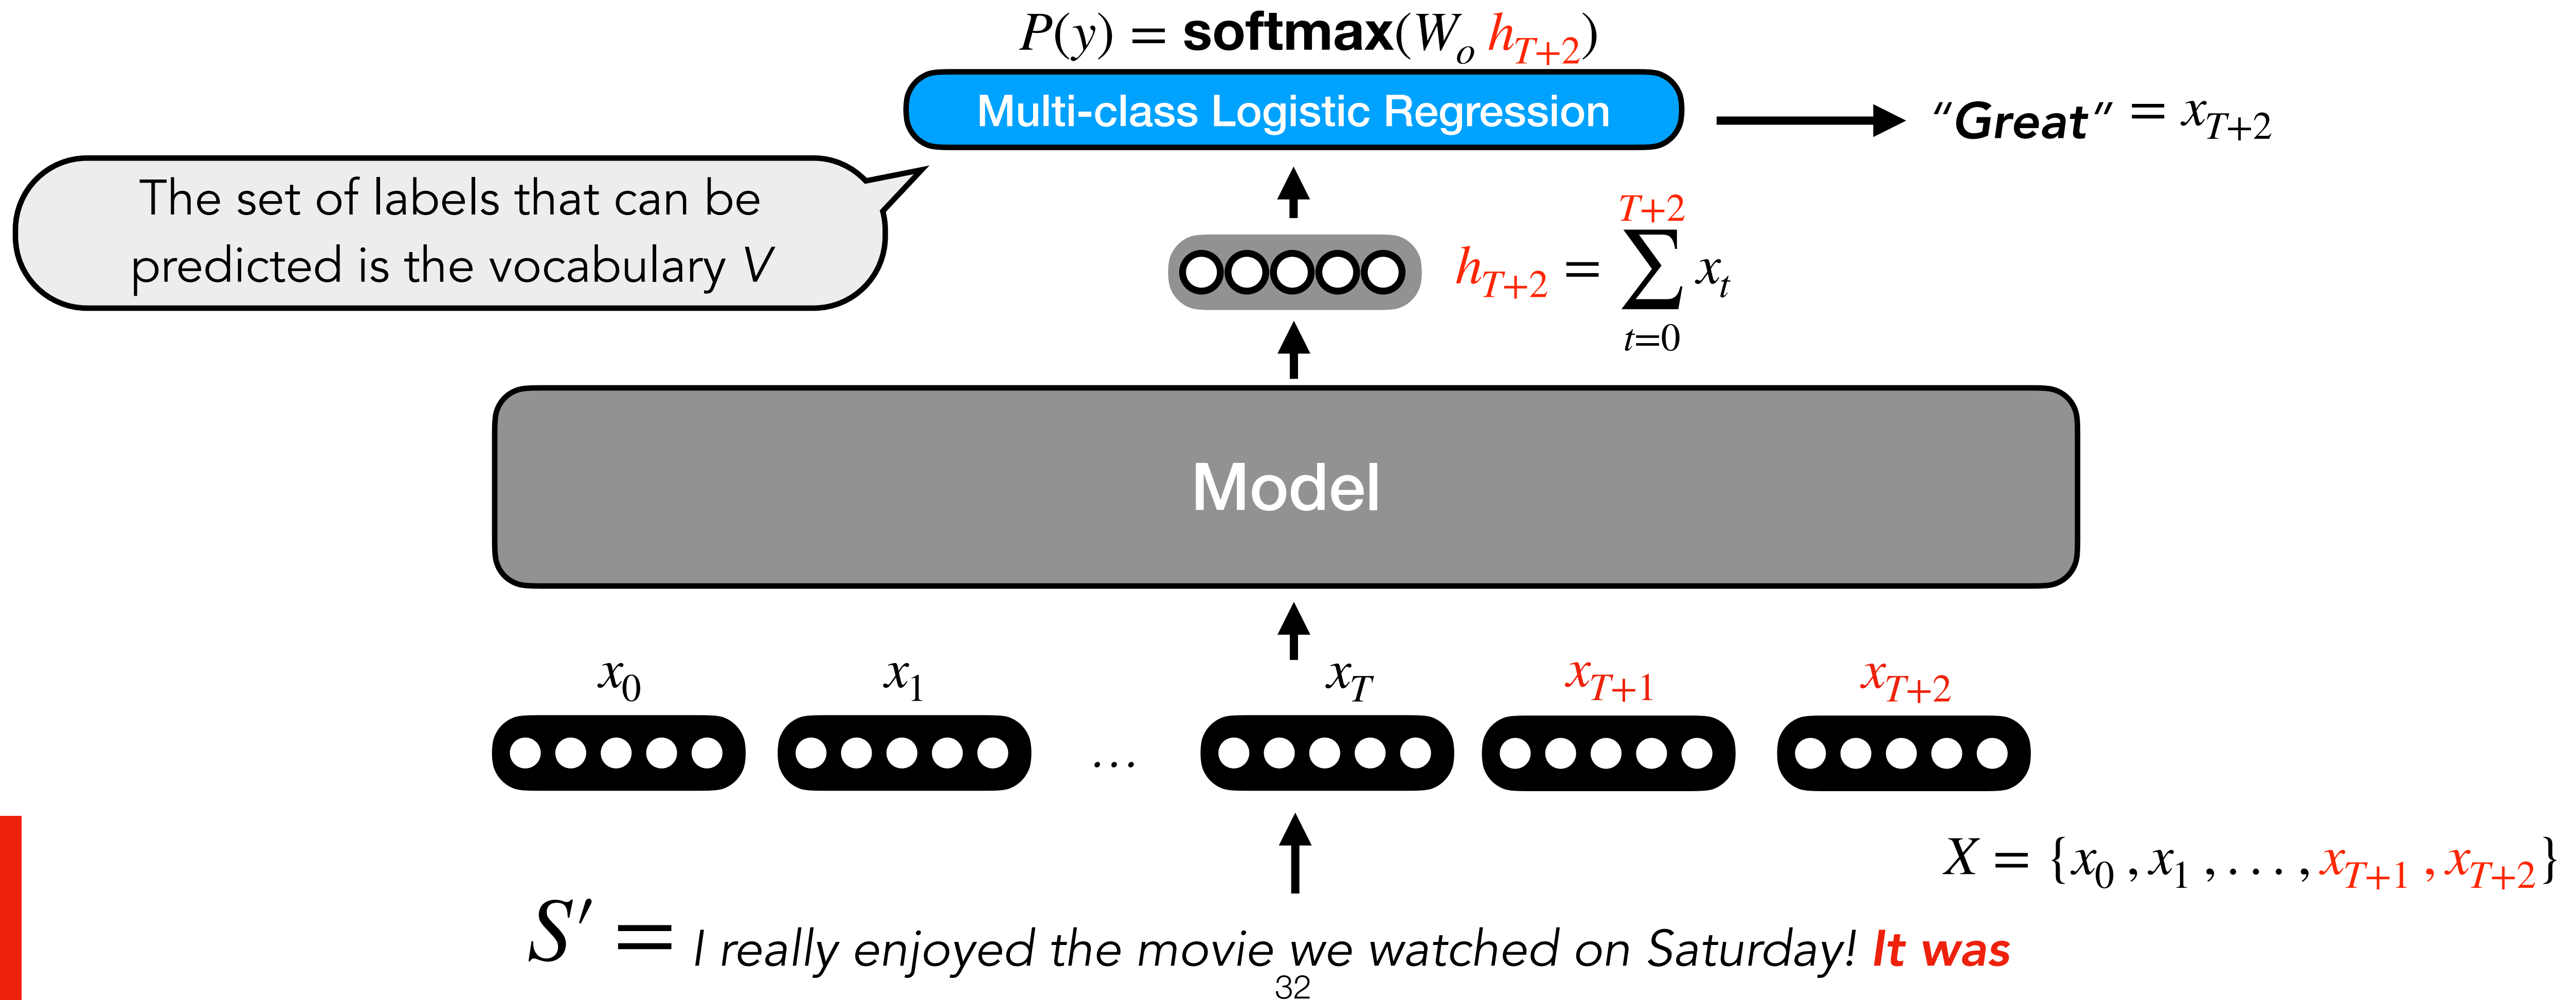
# Text Generation

- **Example:** Generate the next sentence in the review. **Word-by-word!**



# Text Generation

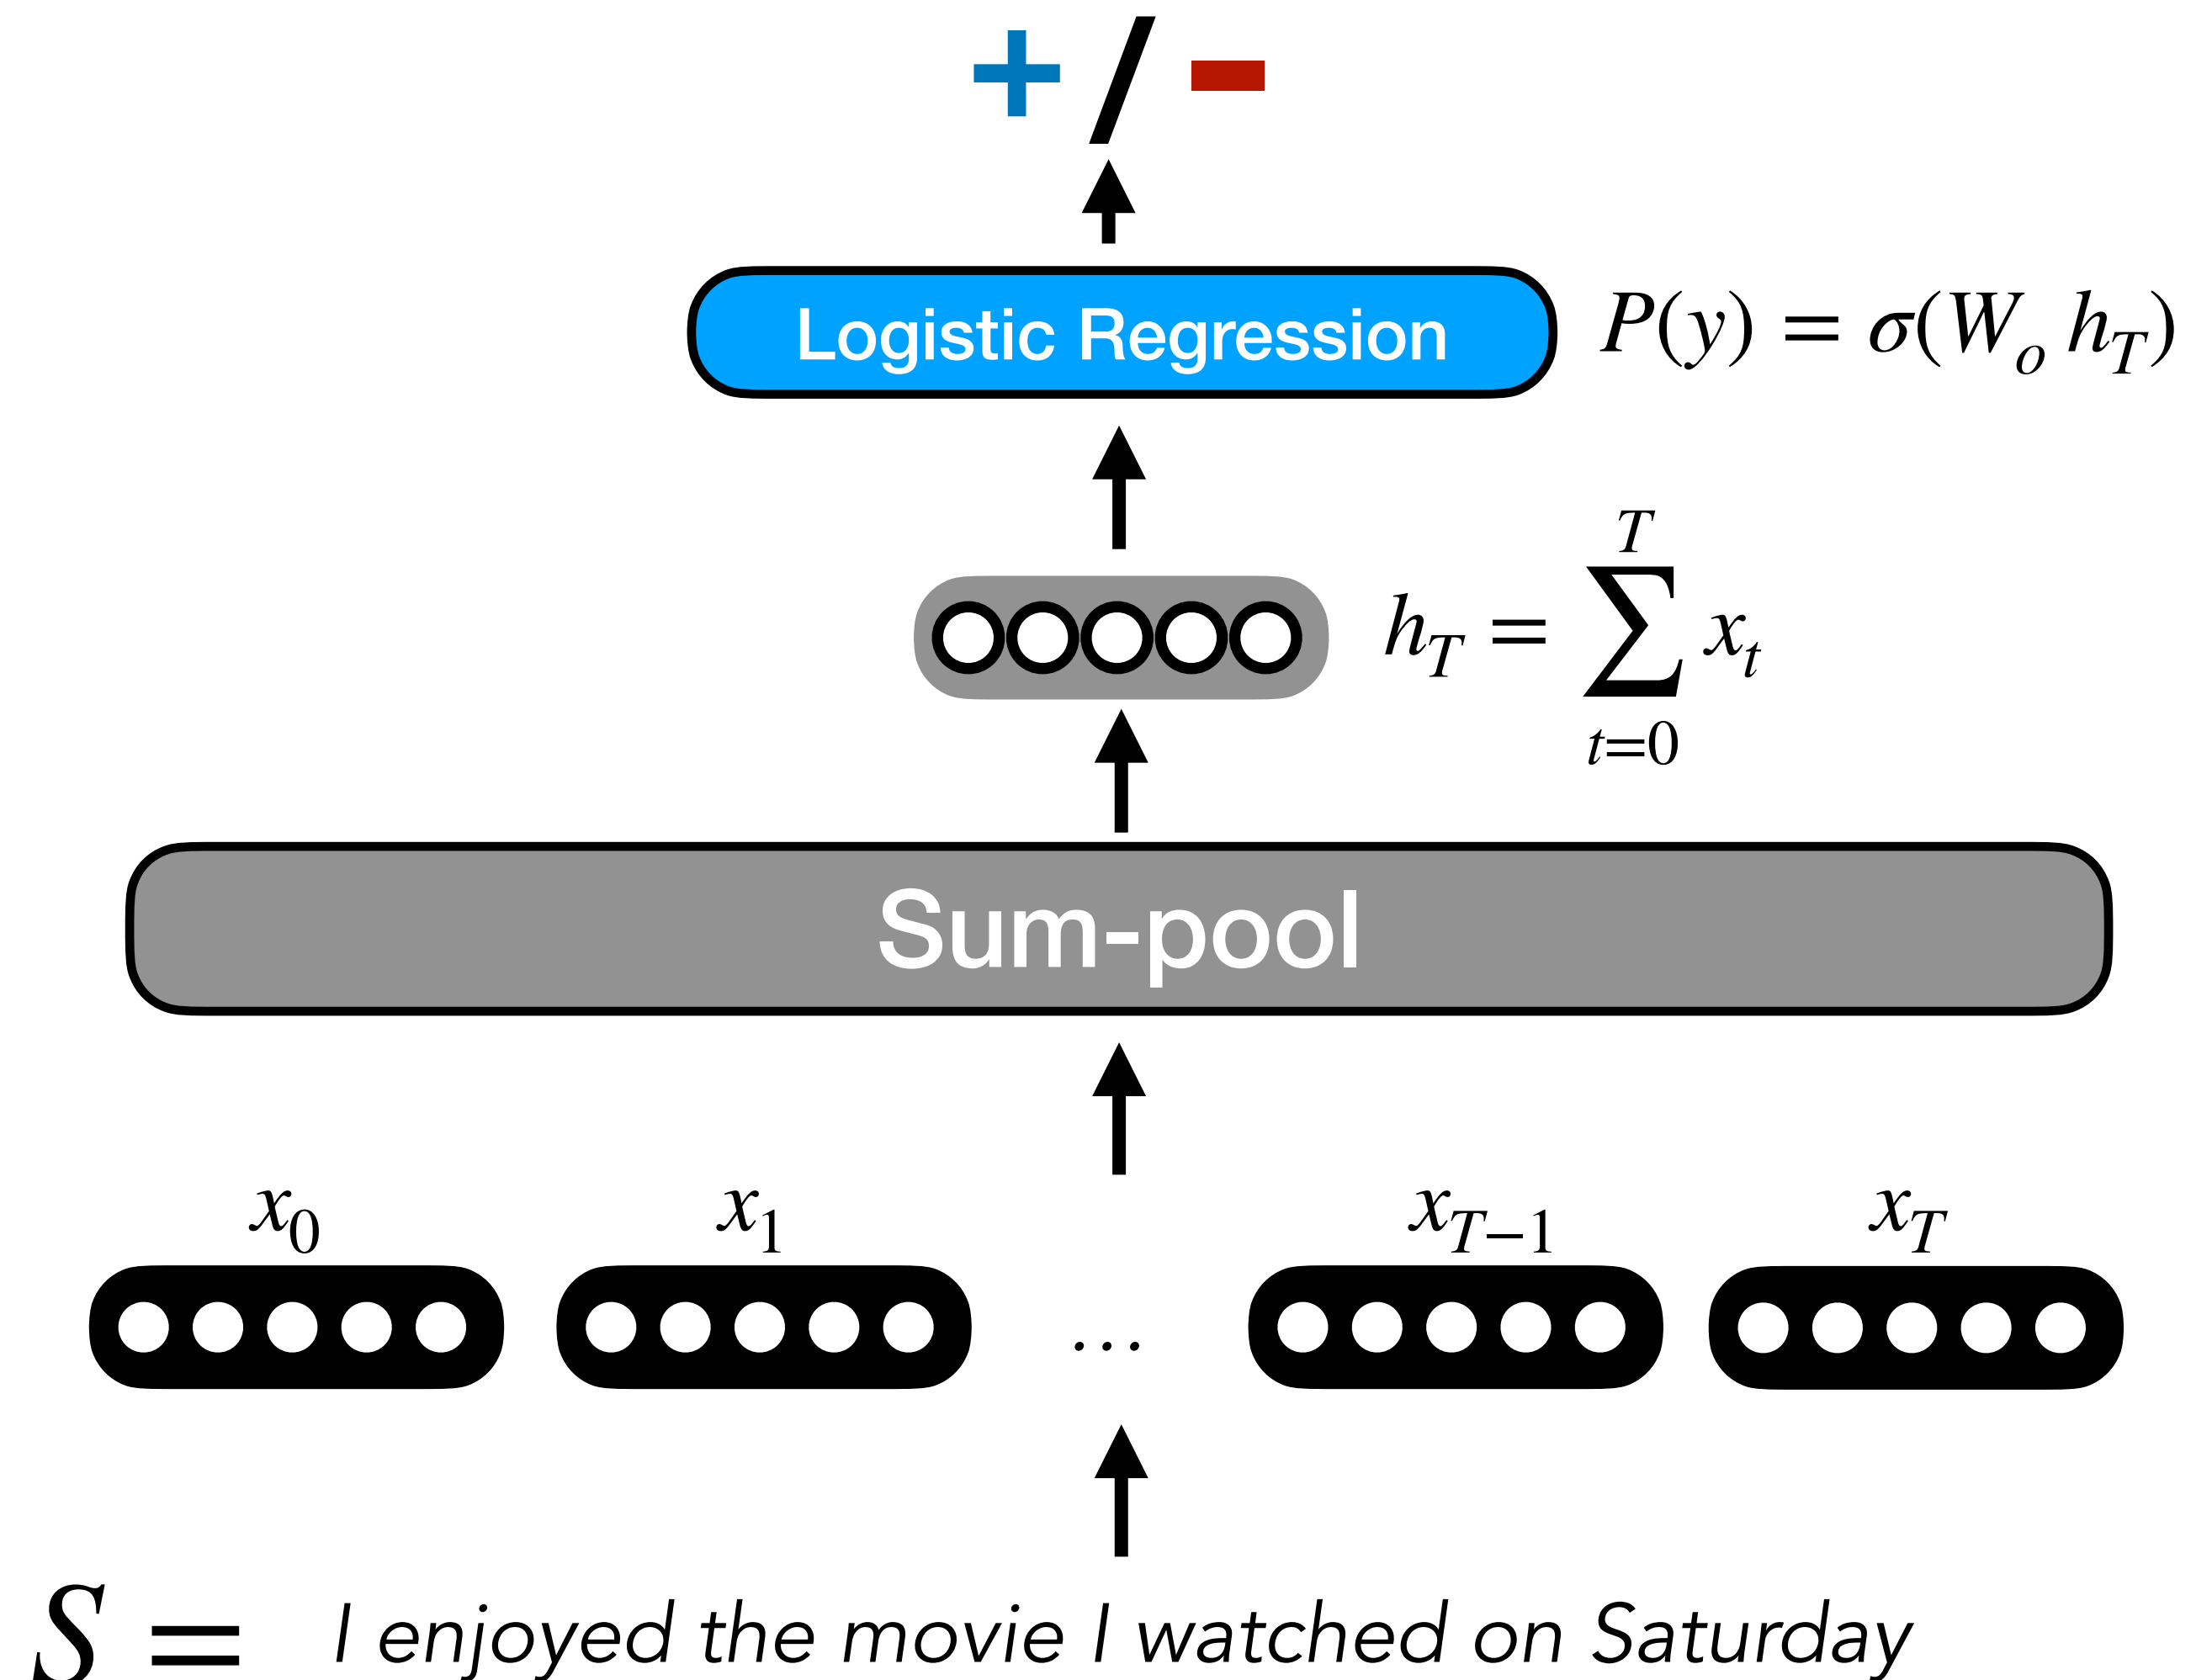
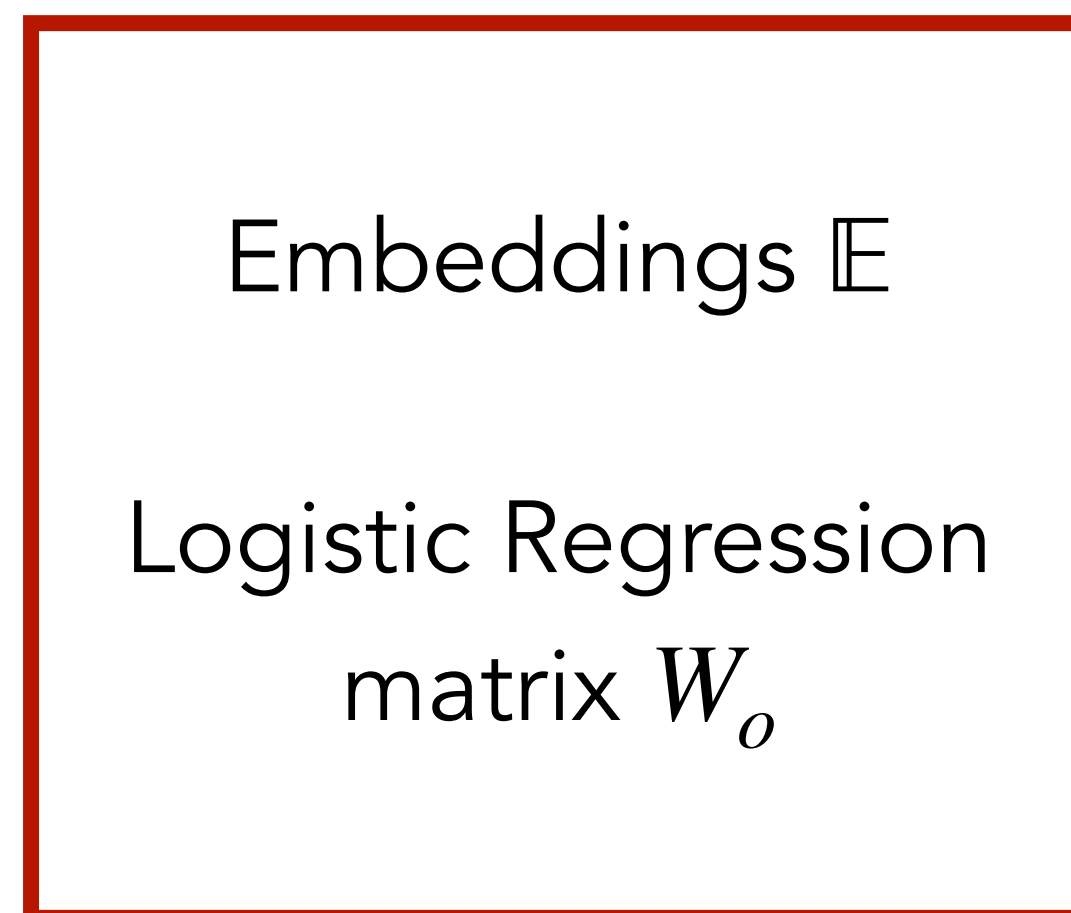
- **Example:** Generate the next sentence in the review. **Word-by-word!**





# Comprehension Questions

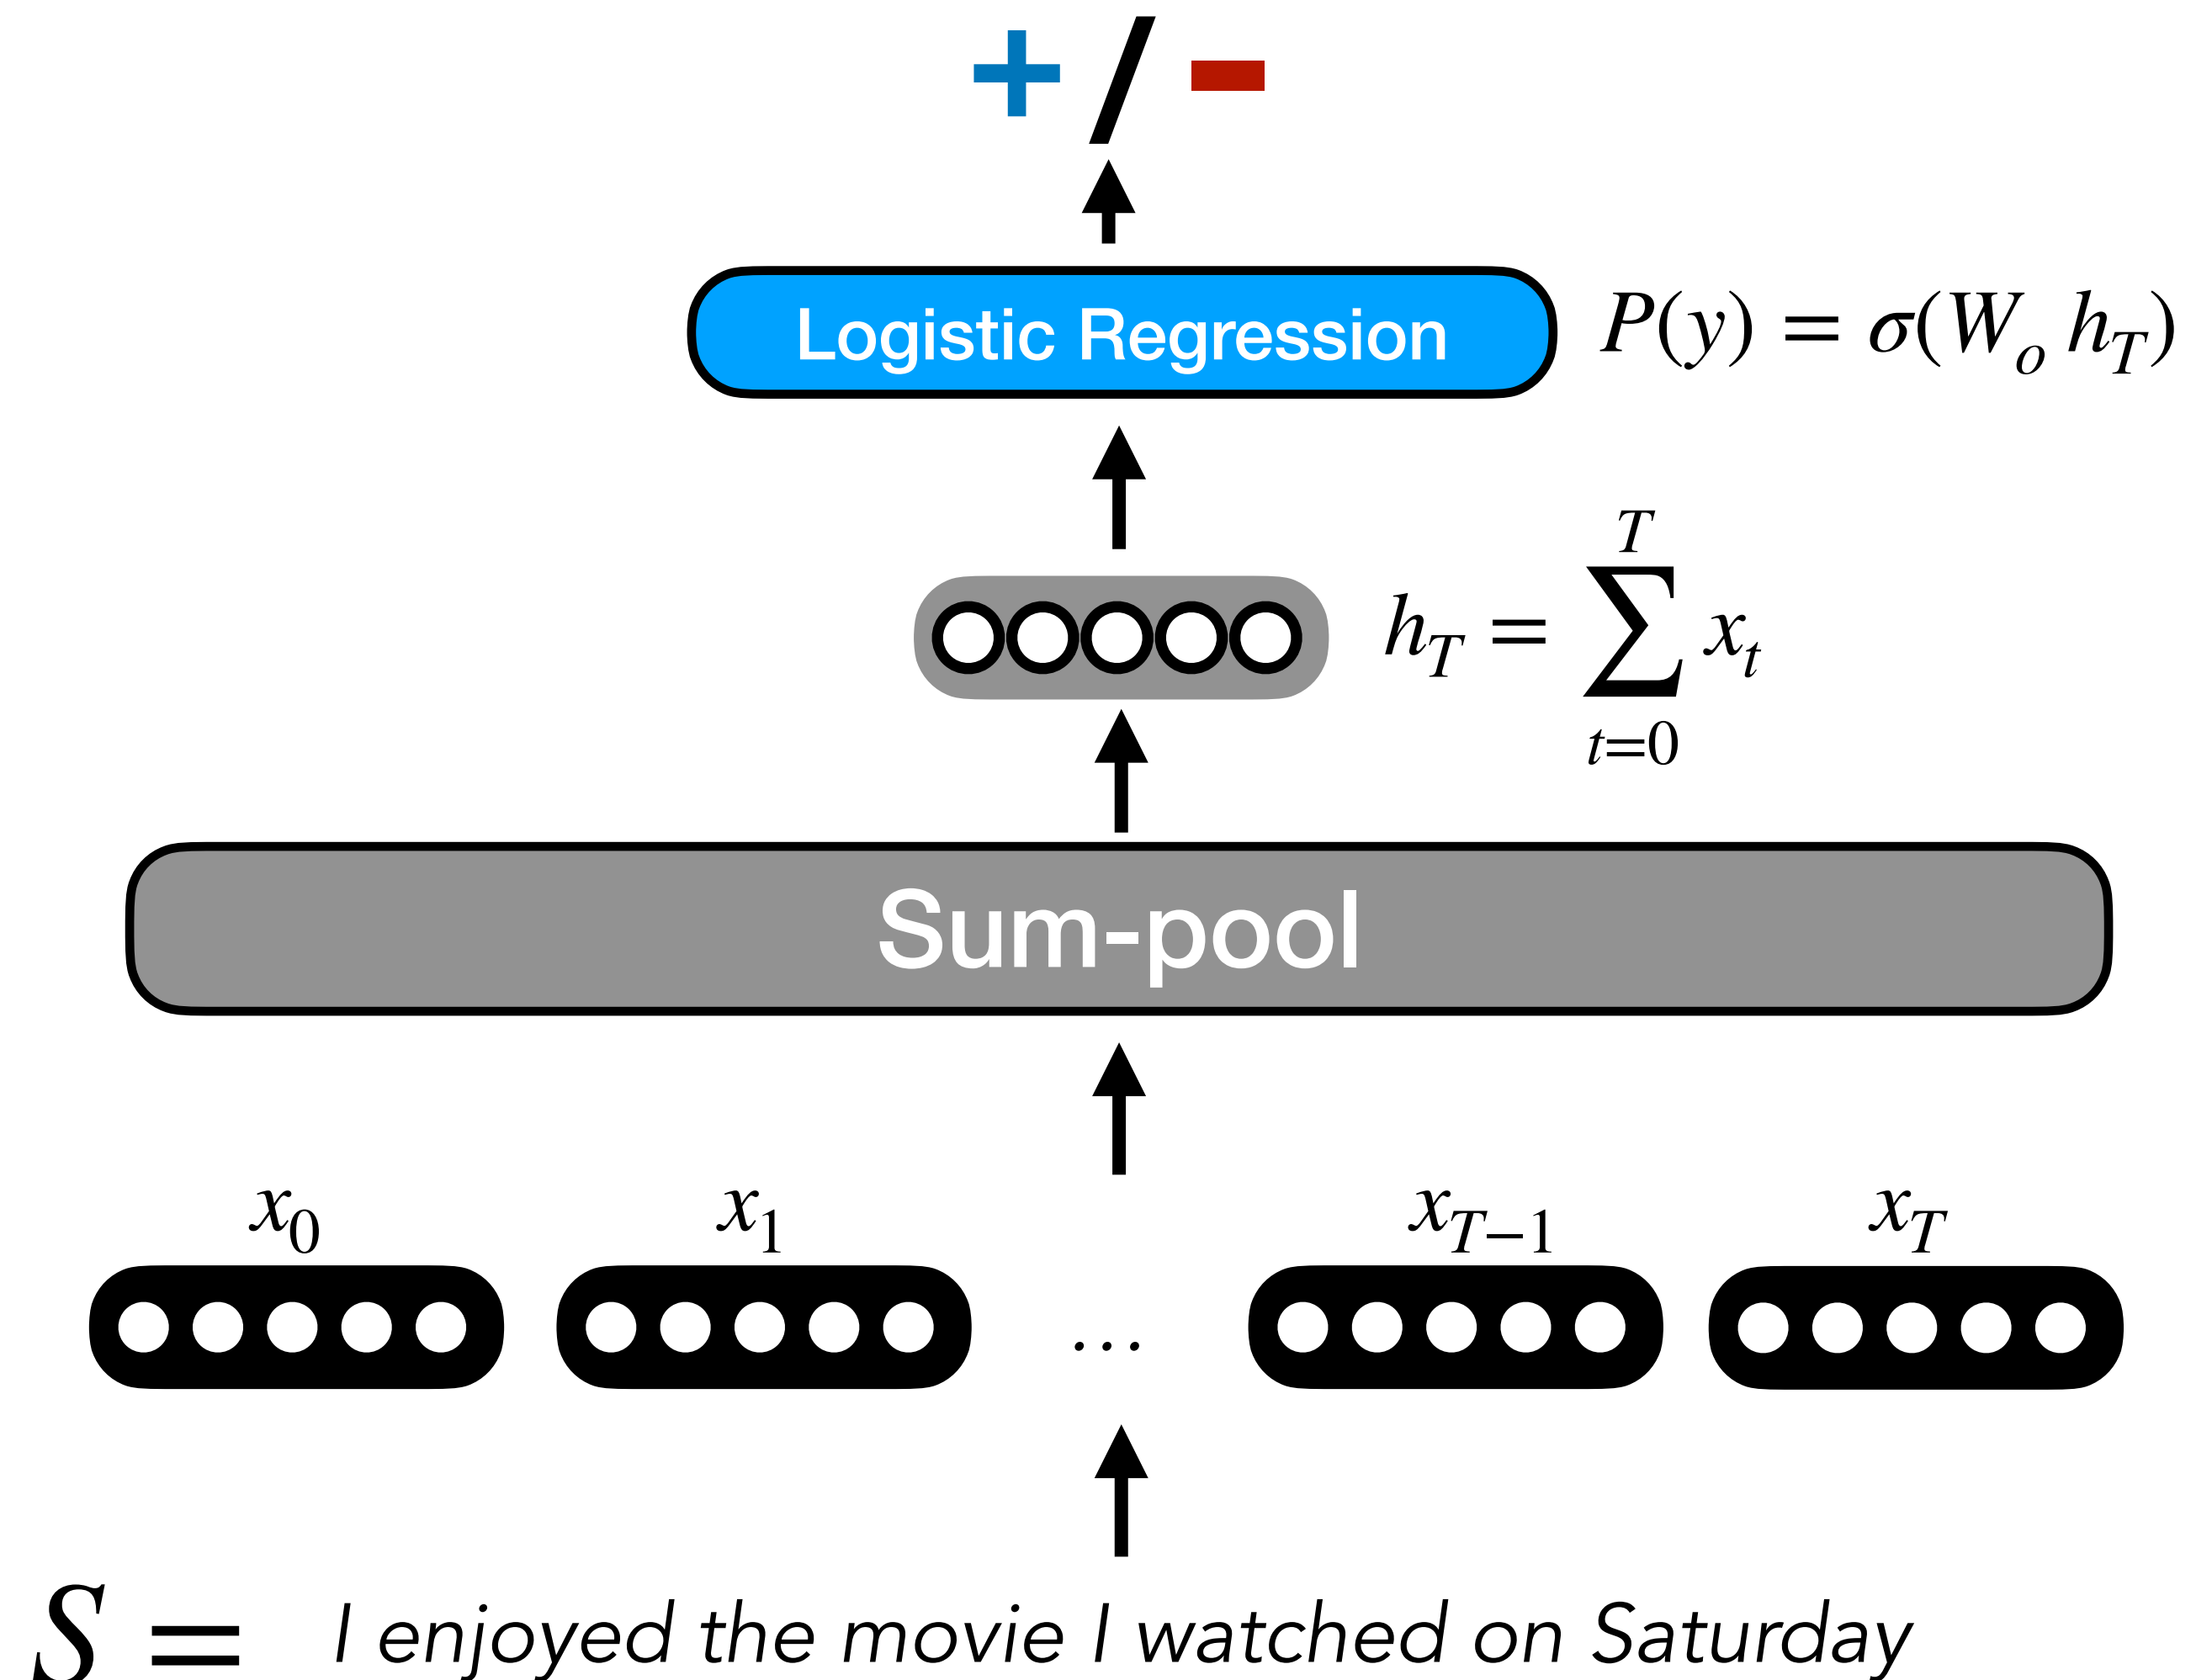
- What are the learnable parameters in our system?



# Comprehension Questions

- What are the learnable parameters in our system?
- How many **unique** embeddings are in  $X$  for this example sentence  $S$ ?

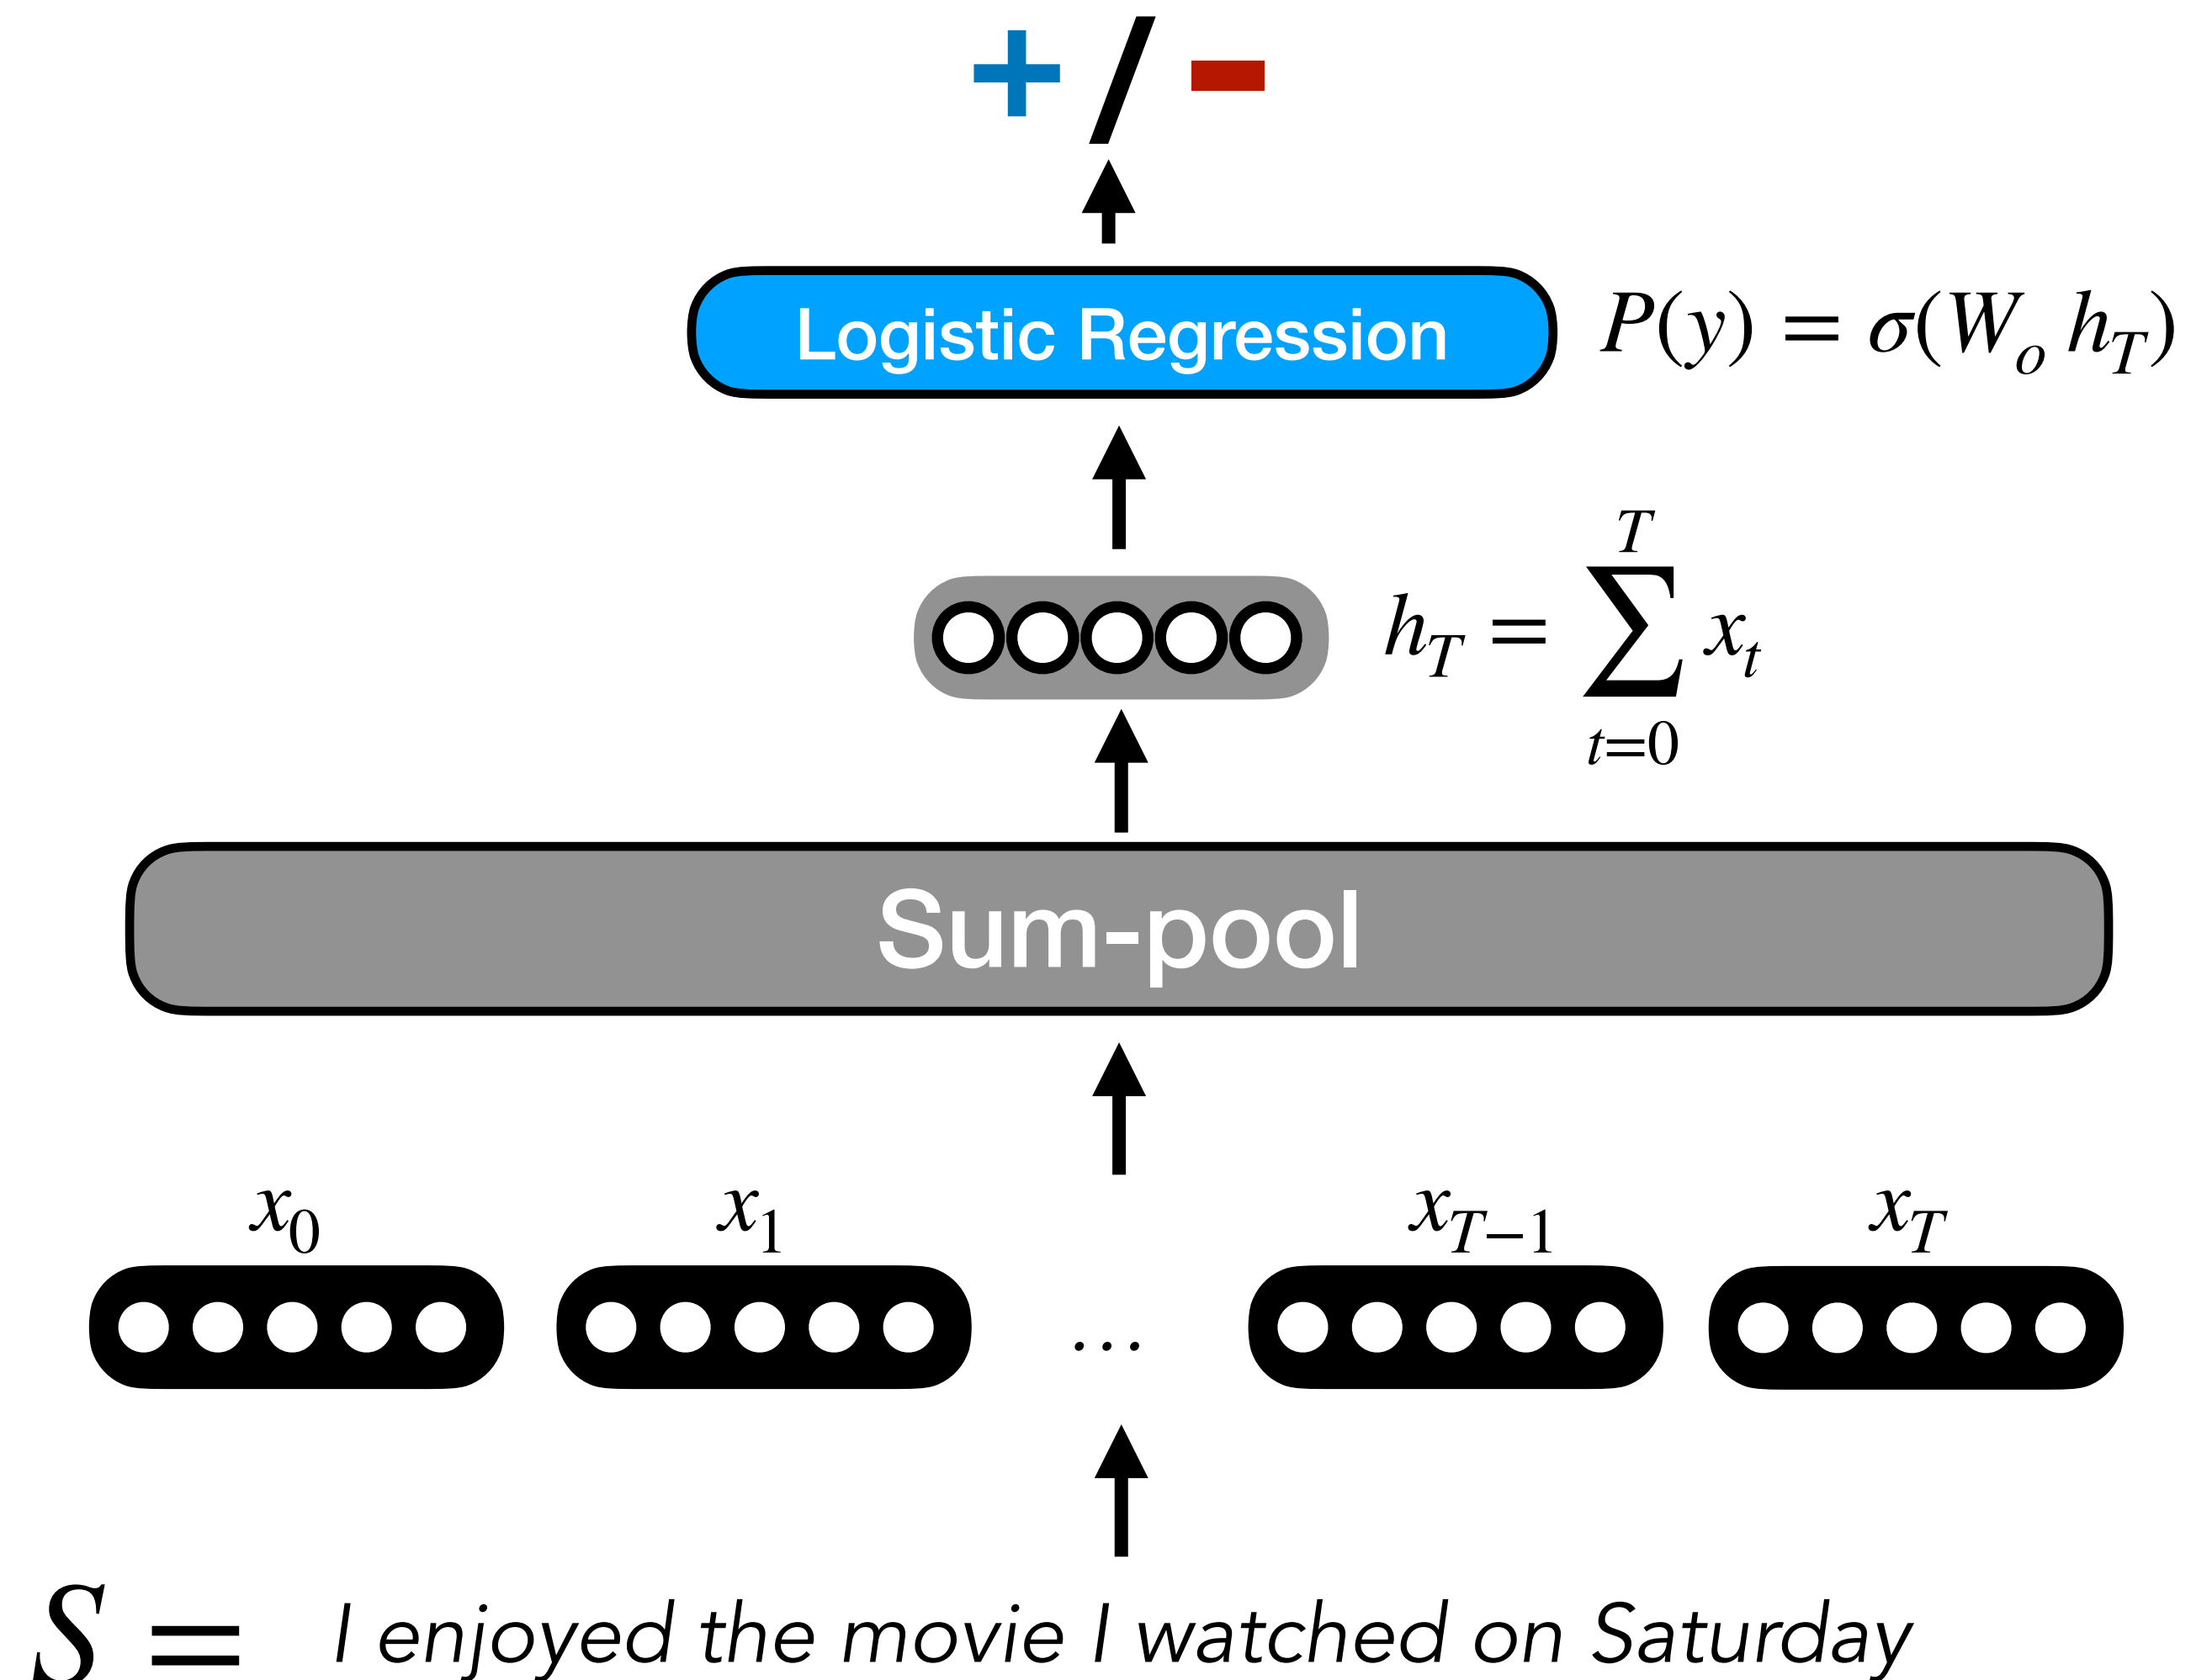
7



# Comprehension Questions

- What are the learnable parameters in our system?
- How many **unique** embeddings are in  $X$  for this example sentence  $S$
- How many **unique** embeddings are in  $\mathbb{E}$  ?

Vocabulary size  $V$



# Recap

- **Words and other tokens become vectors; no longer discrete symbols!**
- Define a vocabulary of words (or token types)  $V$  that our system can assign to a vector
- Define a model that composes these vectors (or embeddings) of words into some sequence representation
- A classifier can map this representation to a set of labels to make a prediction
- The prediction depends on the natural language task we are trying to accomplish
- By learning to make these predictions, we learn better embeddings for the words in the sequences

# Tomorrow

**What could be a better way to learn word embeddings?**

**Self-supervised learning of word embeddings**

# References

- Shen, D., Wang, G., Wang, W., Min, M., Su, Q., Zhang, Y., Li, C., Henao, R., & Carin, L. (2018). Baseline Needs More Love: On Simple Word-Embedding-Based Models and Associated Pooling Mechanisms. *Annual Meeting of the Association for Computational Linguistics*.